Vulnerability Prioritization

Hhmmm…

1. Analyze vulnerability
   - Assess injection, impact, existence of exploits, …
   - Assess **contextual risk** for each instance
2. **Choose** action on each instance
3. **Take** action on each instance

- **Hundreds** of **new** vulnerabilities **each week**
- How to cope with **all** of them?

Vulnerability Management in Organizations (REMIND)

- Vulnerability **Management**:
  - …

- Fundamental components
  - **Asset Management**
    - Which systems
    - Who is in charge
  - **Vulnerability Prioritization**
    - How to allocate defensive efforts to CVEs

Basic Fact #1 (REMIND)

- **Very few CVEs are actually exploited**
- Just to have an idea: **≈5%** of all CVEs **(!)**

Focus on those!

Basic Fact #2 (REMIND)

- CVSS is **not** a good predictor of which vulnerabilities will be **actually exploited**
- Predicting which vulnerabilities will be exploited is a **huge open problem**

- Every predictor you can think of turns out to be
  - "Low" precision:
    You worry about many vulns **unnecessarily**
  - "Low" recall:
    You wrongly **neglect** many vulns

Our path

- Exploit Prediction Scoring System (**EPSS**)

- Considered the "state of the art"
- ...but it does have many limitations

- We will see:
  - How it works
  - Public data on its assessment

  - Some of its limitations

Exploit Prediction: Problem Definition

Exploit Probability: Definition (I)

- Vulnerability `CVE-i`
- `P(CVE-i, d)`:

  $$\frac{\#\text{ Exploitation attempts of CVE-i in [d,d+30]}}{\#\text{ Expolitation attempts of all CVEs in [d,d+30]}}$$

- Probability that `CVE-i` will be exploited in the next 30 days

- **Against whom?**
- **How computed?**

Exploit Probability: Definition (II)

- Vulnerability `CVE-i`
- `P(CVE-i, d):`

  `# Exploitation` **attempts** of `CVE-i` in `[d,d+30]`

  `# Expolitation` **attempts** of **all** `CVEs` in `[d,d+30]`

- Probability that `CVE-i` will be exploited in the next 30 days

- **Worldwide** (everywhere)
- **Approximated** by collecting many TI feeds
- Computed **a posteriori**

Exploit Prediction: Problem Definition (I)

- We define a criterion for choosing **which vulnerabilities to patch**
- Subset of **all known vulns** at t_0

Many possible criteria

- All vulns with CVSS >= 9 (Critical)
- All vulns with remote injection
- All vulns of Windows software
- ...

Remark

- Subset of **all known vulns** at t_0

- An organization should focus **only** on vulns on **its** systems (and their **risk**)
- We are **pretending** all the known vulns are **equally relevant**

Exploit Prediction: Problem Definition (II)

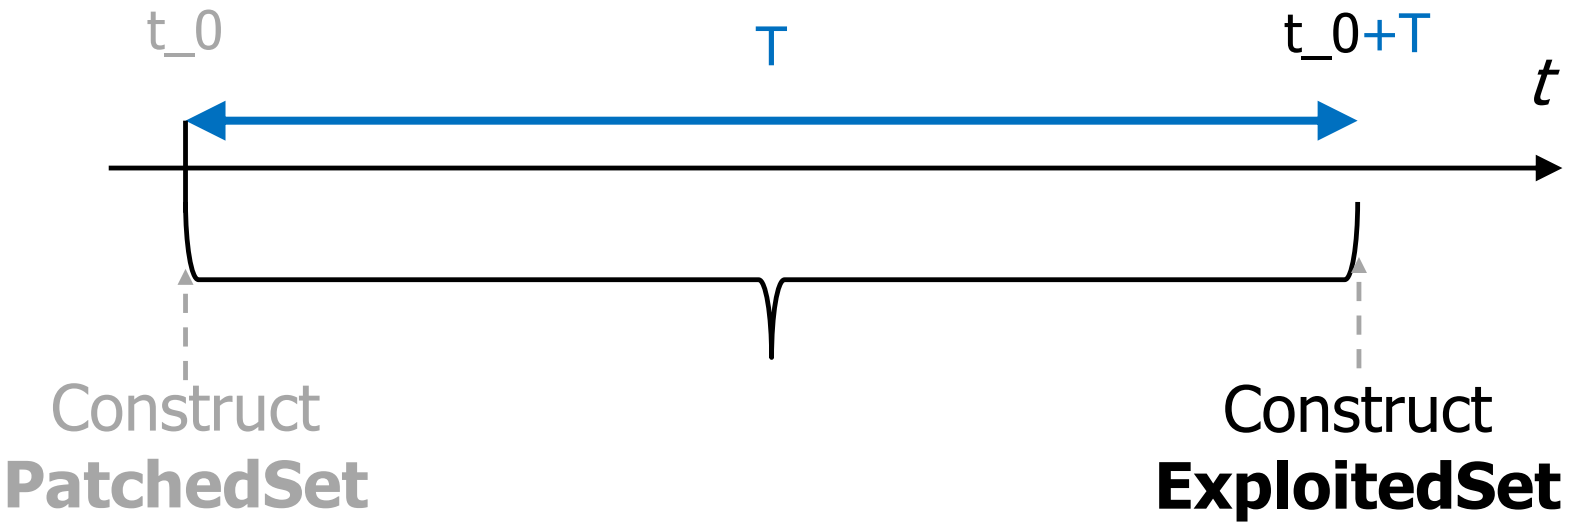

- Vulnerabilities that have been **actually exploited worldwide** in T
- Those with `P(CVE-i)≠0` (on some day in `[t_0,t_0+T]`)

Exploit Prediction: Problem Definition (III)

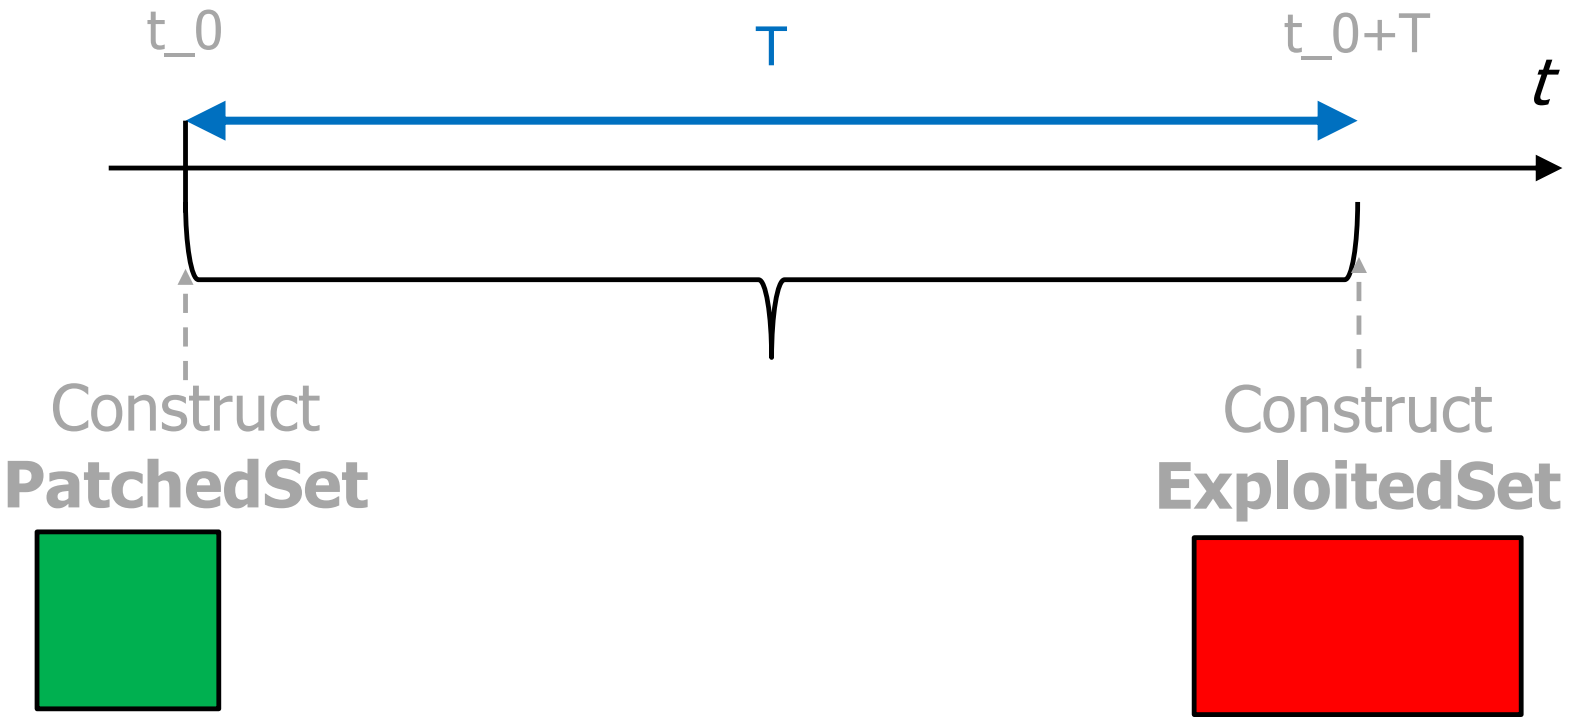

- We "compare" PatchedSet vs ExploitedSet

Efficiency (Precision) (I)

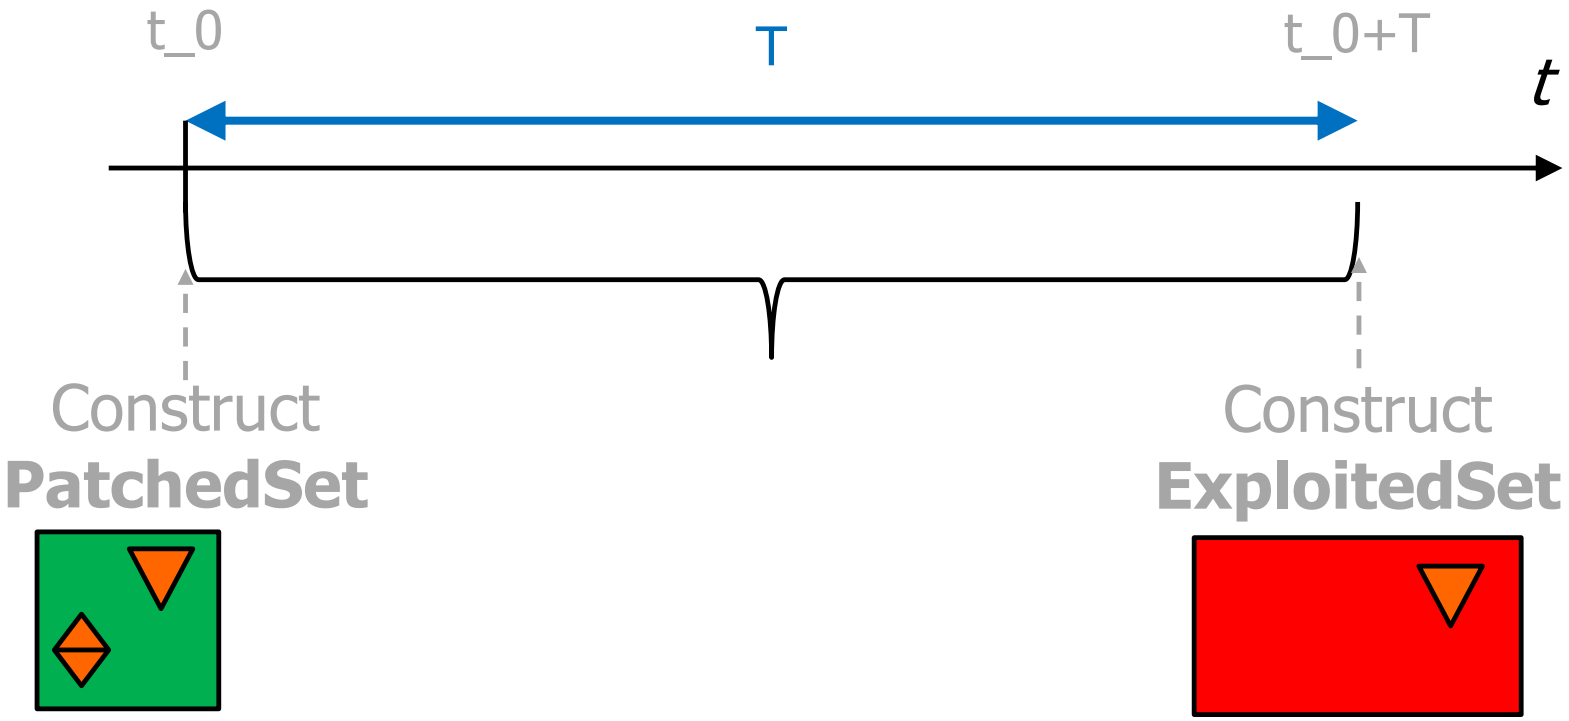

How many Patched vulns will be Exploited?

Efficiency (Precision) (II)

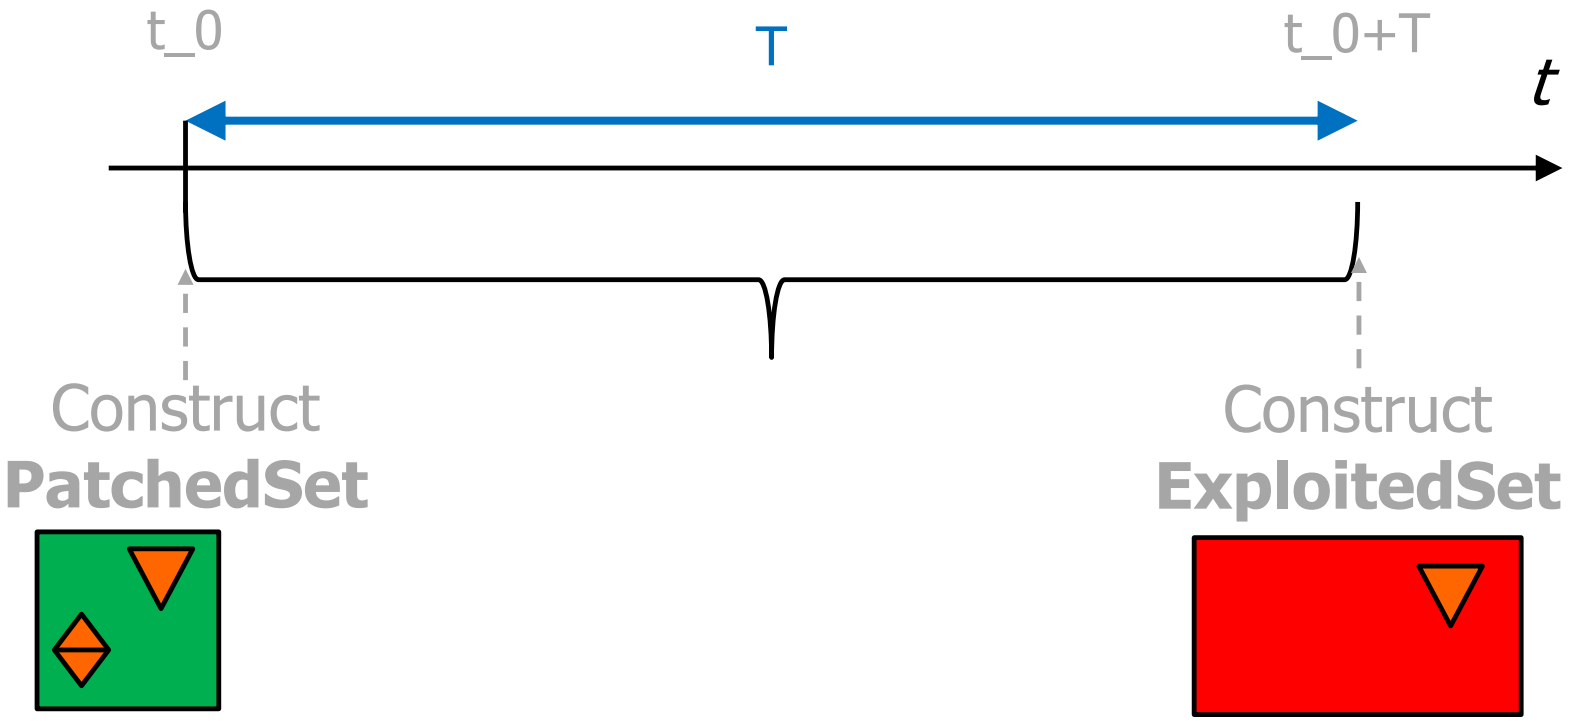

Efficiency

$$\frac{\#(\text{Patched } \textbf{and Exploited})}{\#(\text{Patched})}$$

Coverage (Recall) (I)

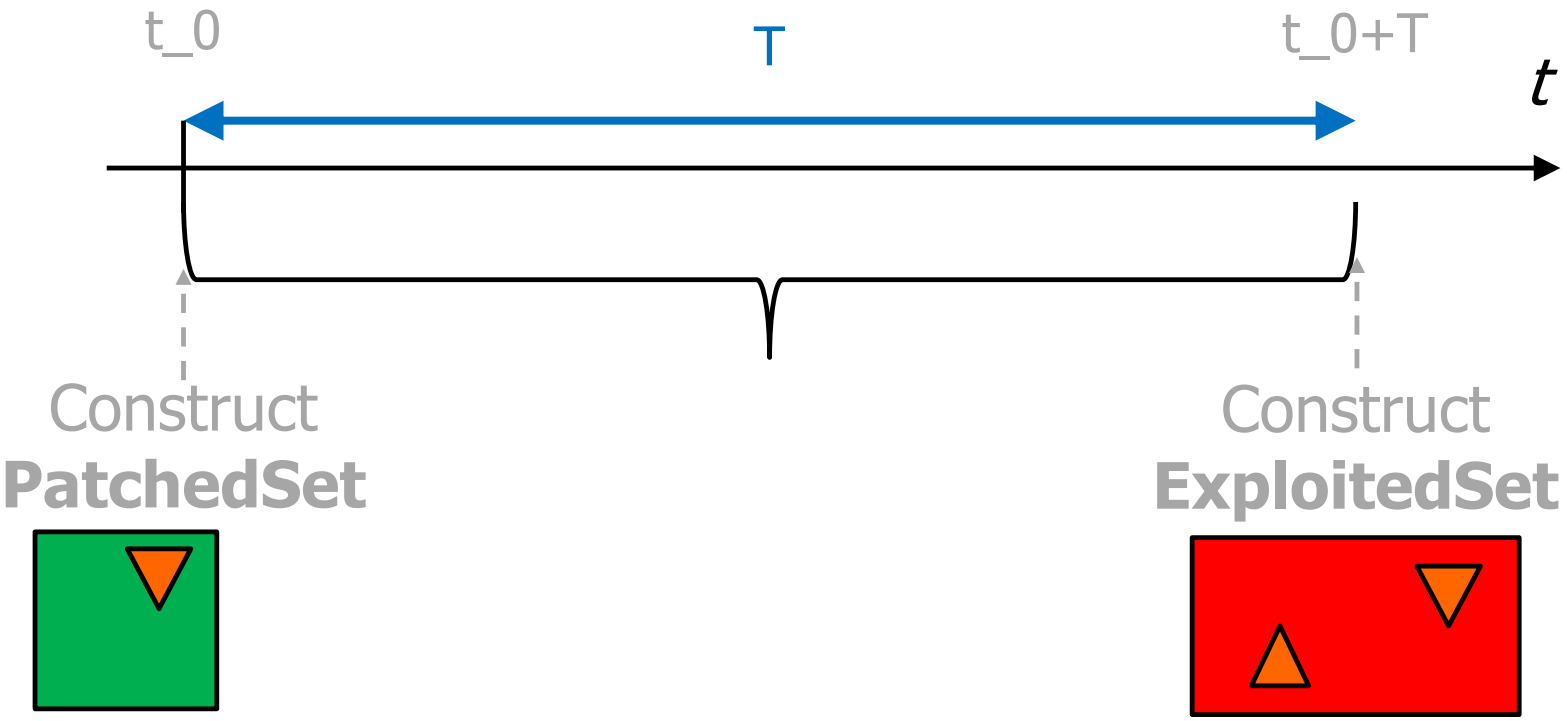

How many Exploited vulns we Patched?

Coverage (Recall) (II)

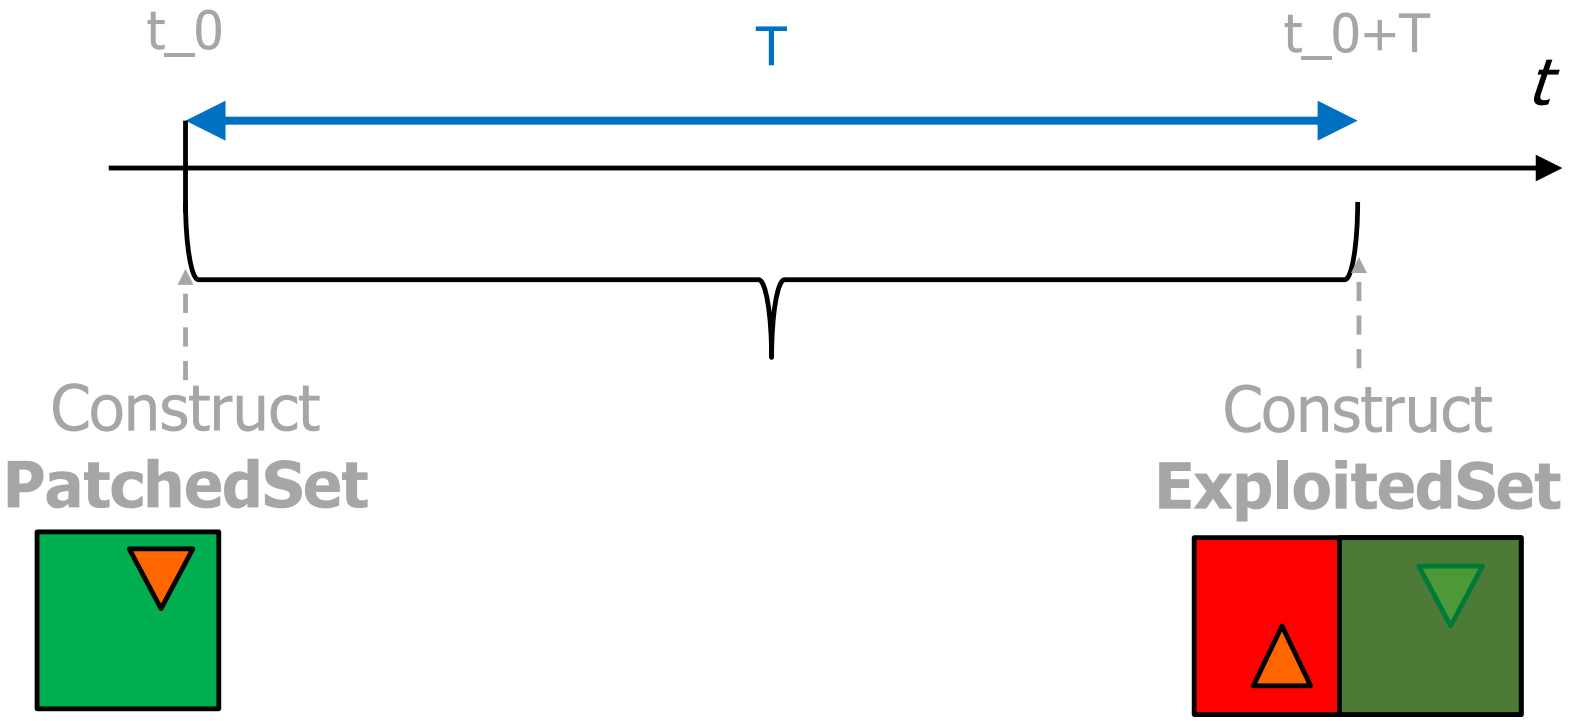

Coverage:

$$\frac{\#(\text{Exploited } \textbf{and Patched})}{\#\ (\text{Exploited})}$$

High Efficiency / Low Coverage

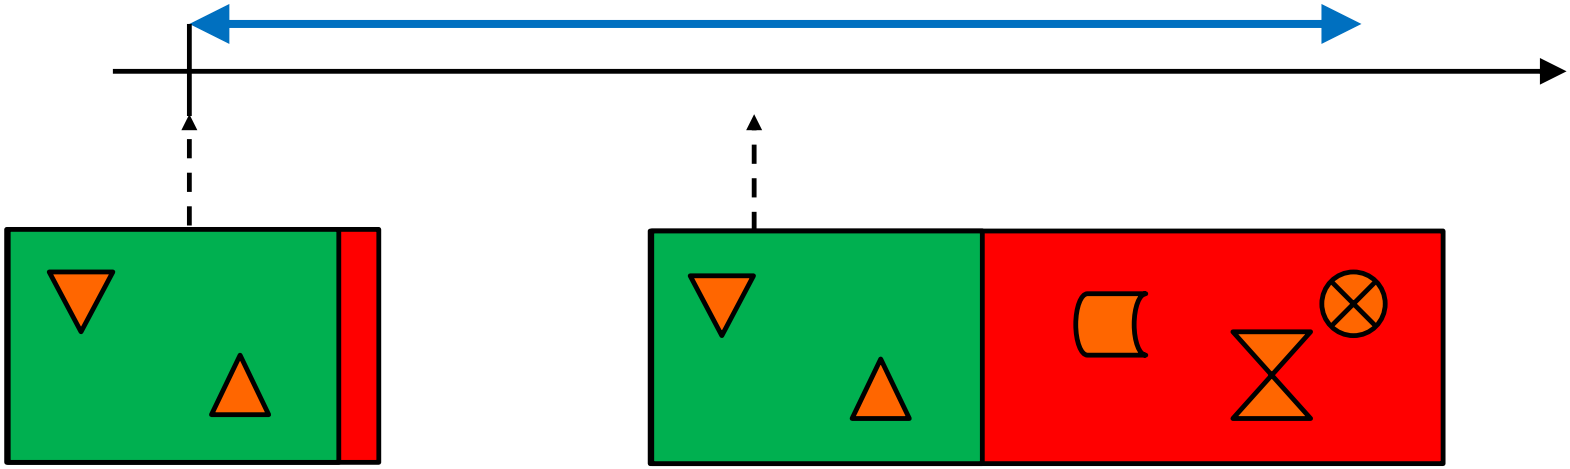

I have patched mostly vulns that matter

...but I have missed many vulns

Low Efficiency / High Coverage

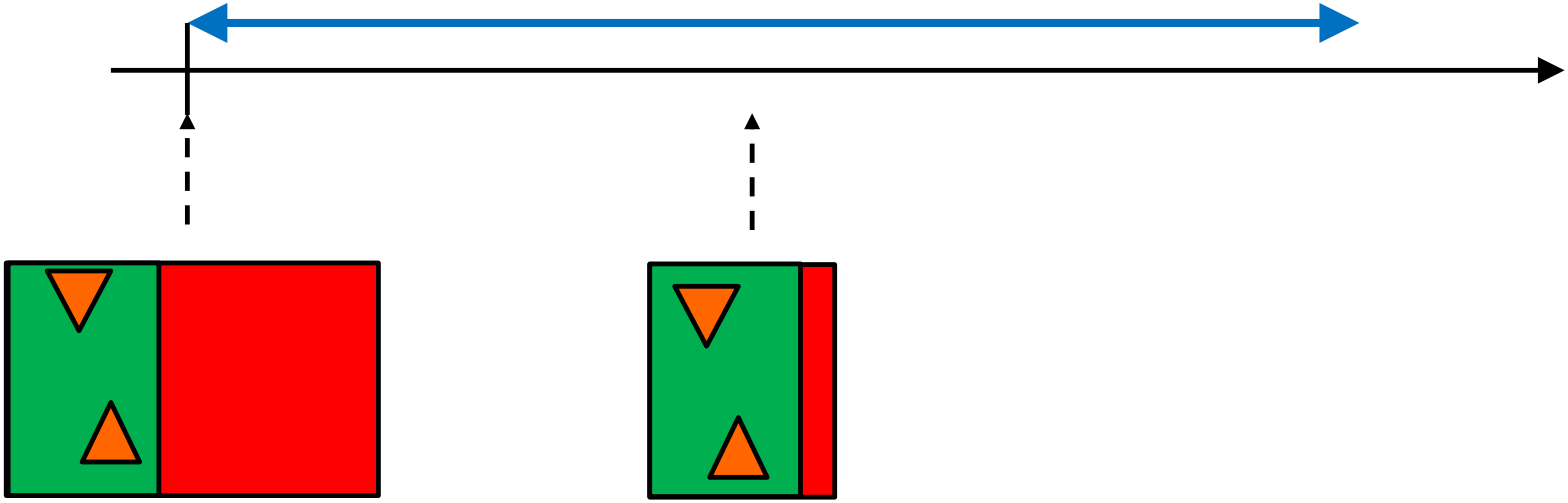

I have wasted lot of patching effort

...but I have covered nearly all vulns that matter

Efficiency

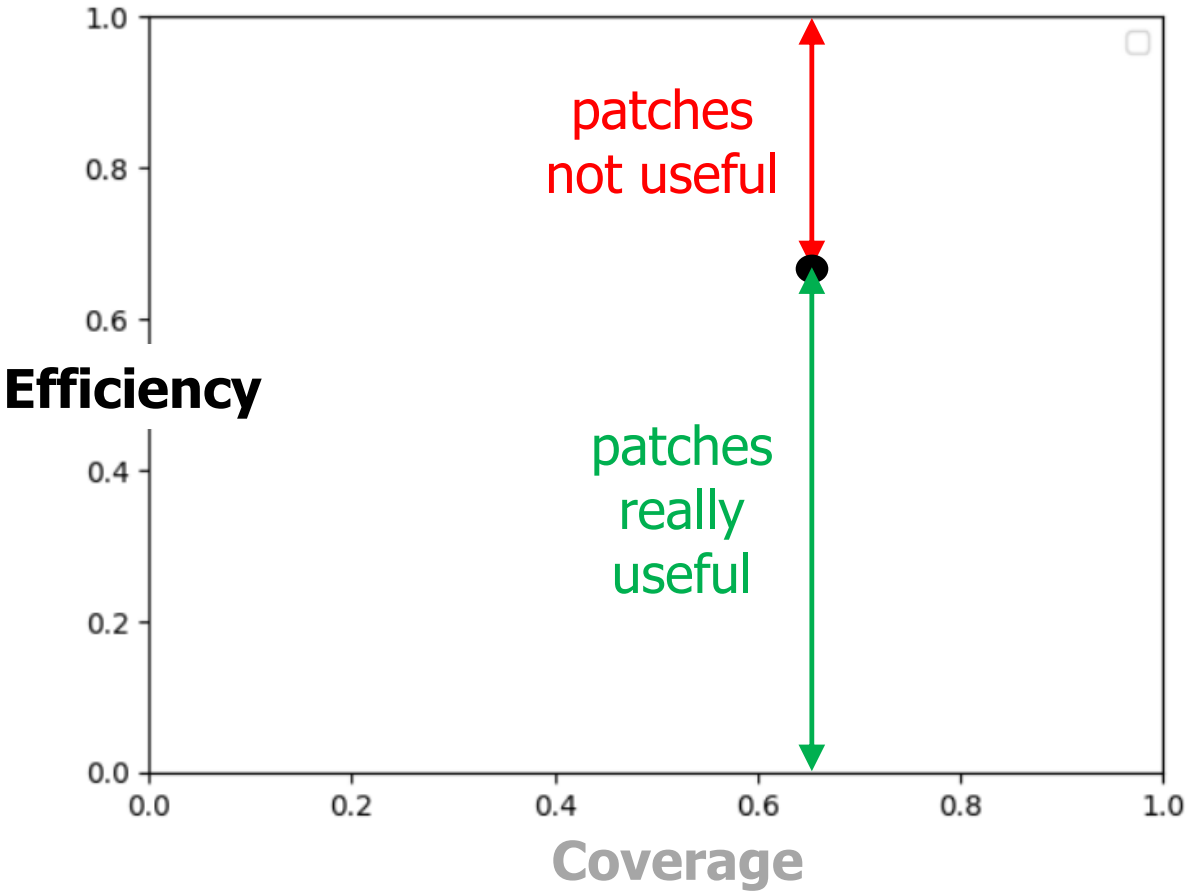
1.0
0.8
0.6
0.4
0.2
0.0
Efficiency
0.0
0.2
0.4
0.6
0.8
1.0
Coverage
patches
not useful
patches
really
useful

Coverage

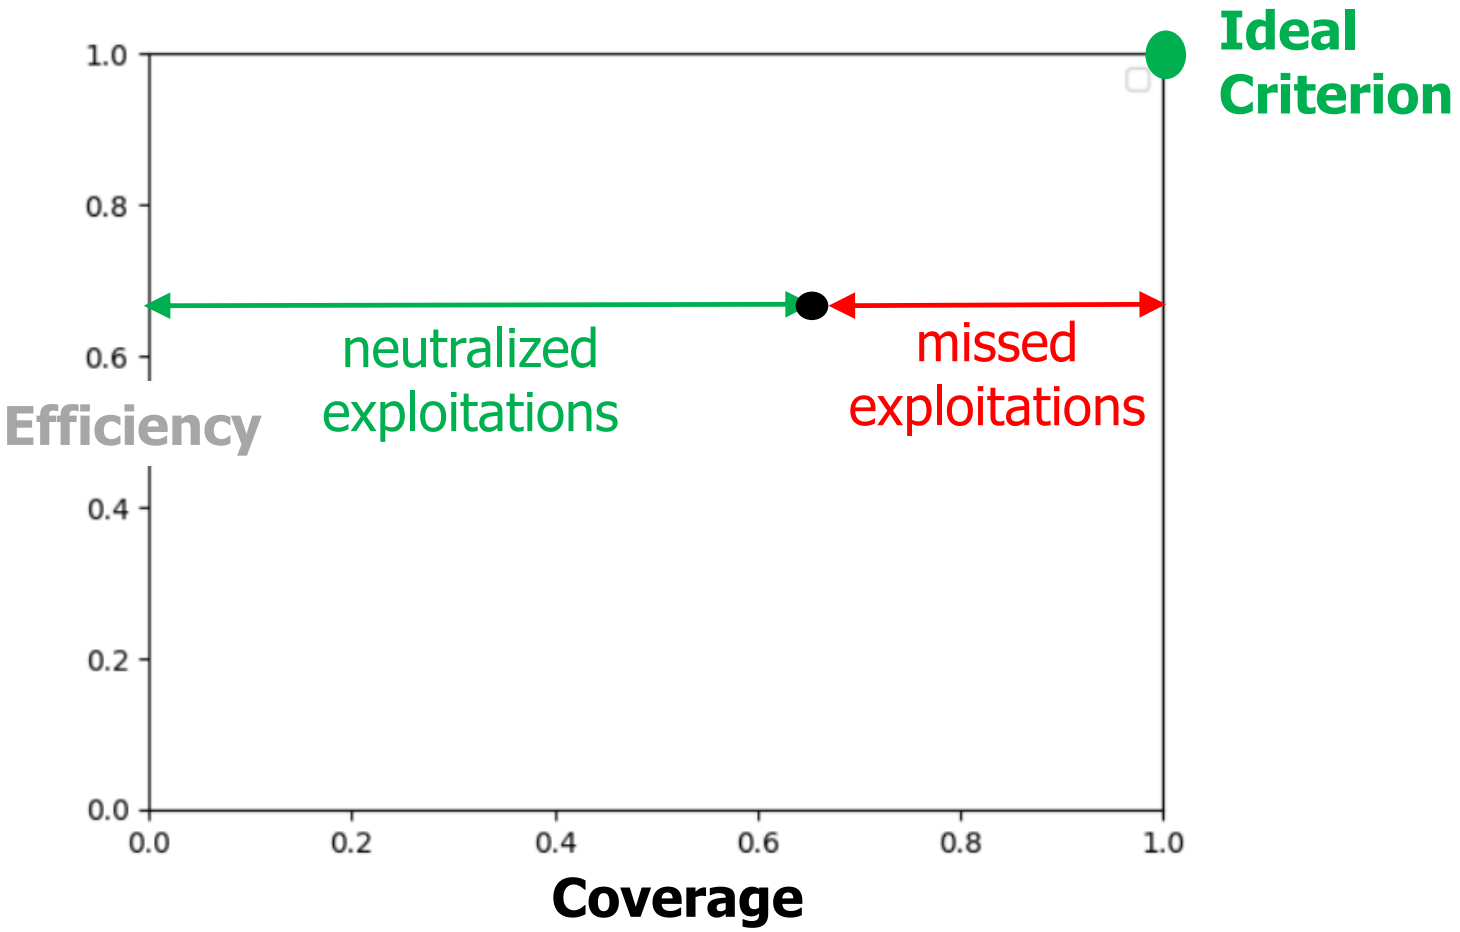
Ideal
Criterion
1.0
0.8
0.6
0.4
0.2
0.0
Efficiency
neutralized
exploitations
missed
exploitations
0.0
0.2
0.4
0.6
0.8
1.0
Coverage

Ideal Criterion

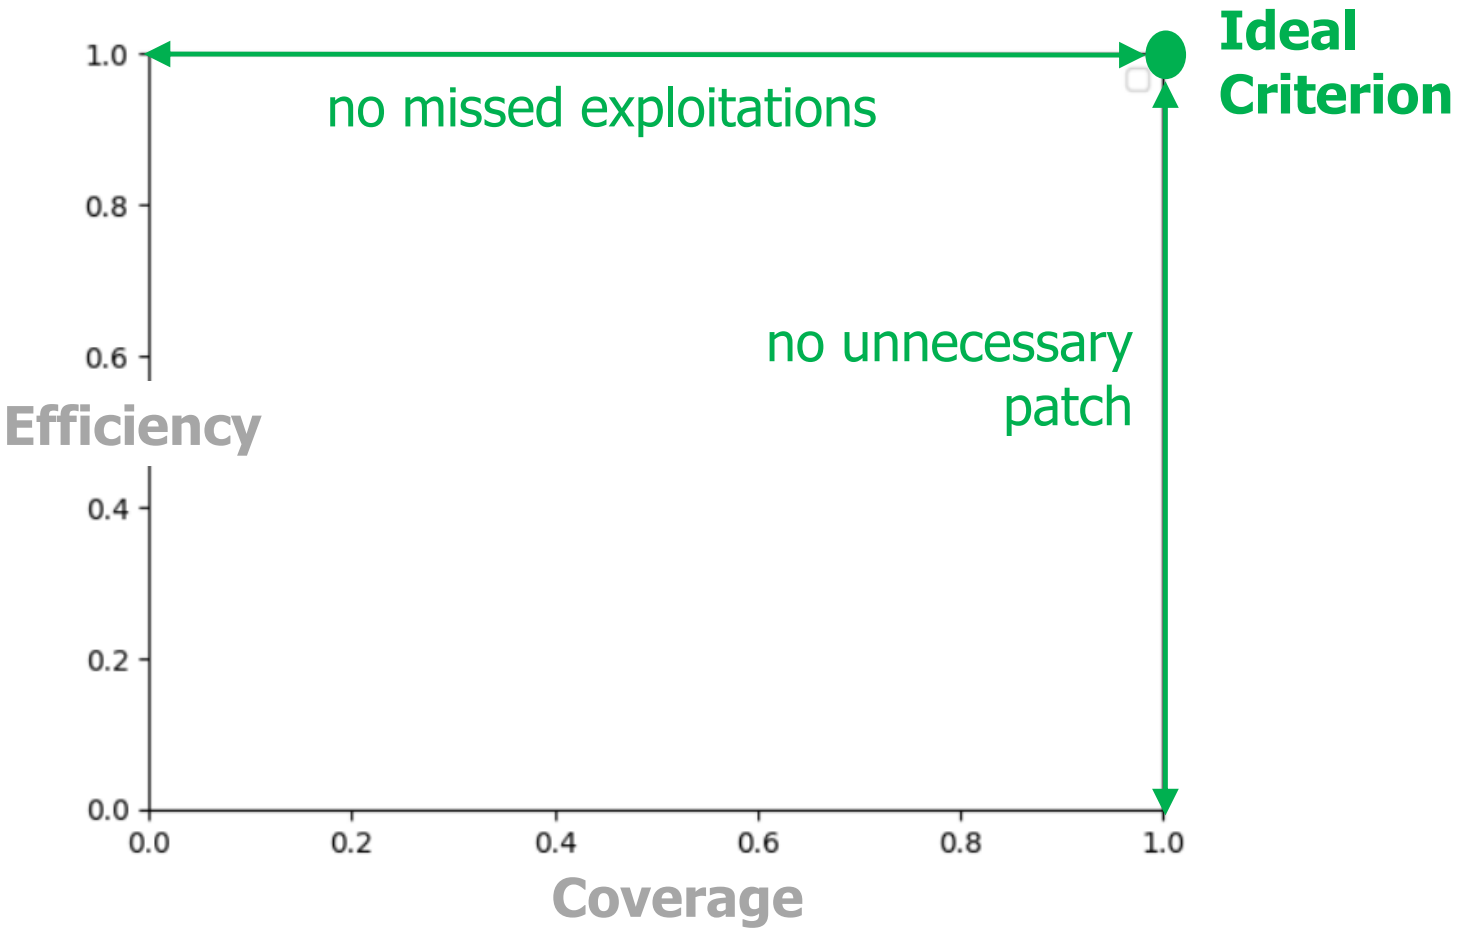

Patching Effort (I)

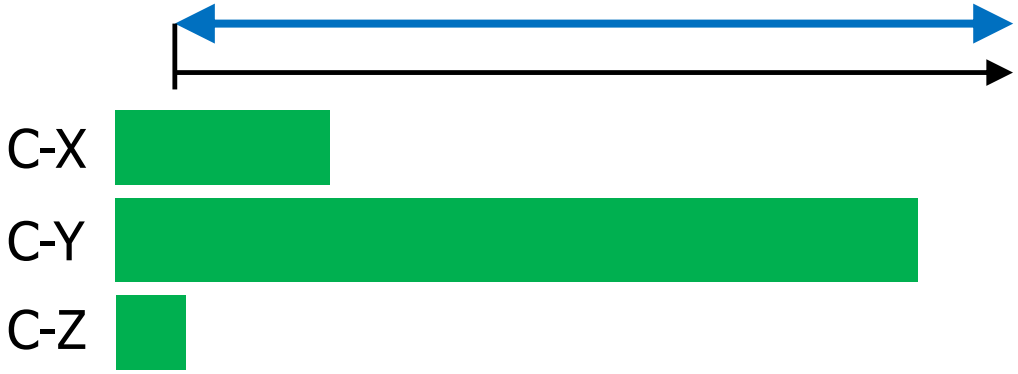

- #(PatchedSet) ≡ **Patching Effort**
- It depends on the chosen criterion:
  - All vulns with CVSS Critical
  - All vulns with remote injection
  - All vulns of Windows software
  - ...

Patching Effort (II)

- Coverage and Efficiency are **relative** indexes
- **Independent** of Patching Effort

- You could have many **different** criteria with:
  - **Identical** Coverage / Efficiency
  - Widely different PatchingEffort

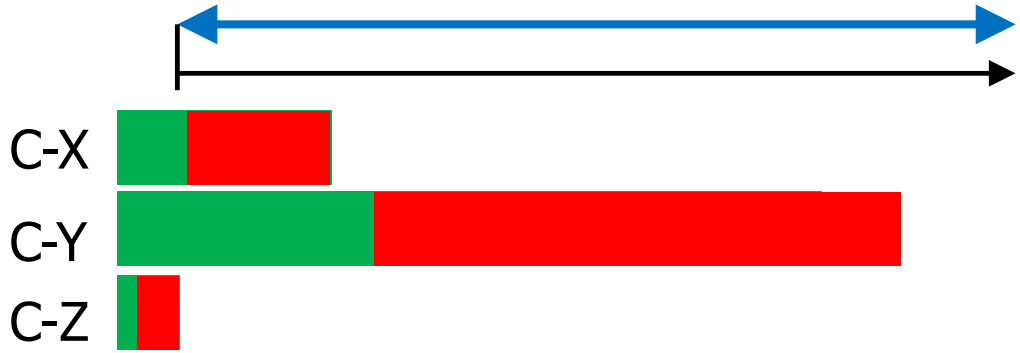

Patching Effort (II)

- Coverage and Efficiency do **not** tell the whole story
- Assessment of any given criterion requires analyzing all **the 3 indexes**

- Example:
  - Criterion C-Y has excellent Coverage / Efficiency
  - …but our org **cannot afford** its PatchingEffort

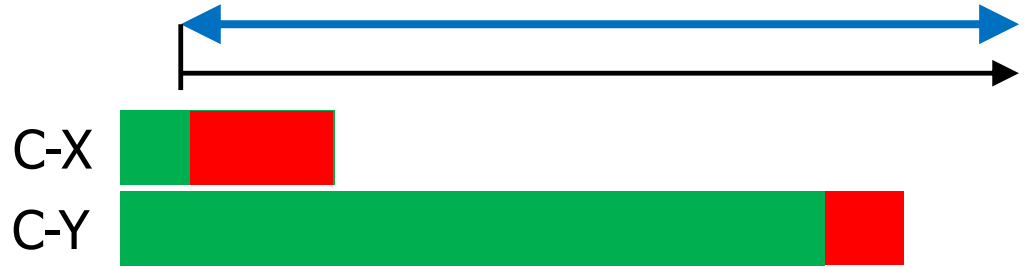

Exploit Prediction: Problem Definition Summary

- Criterion for **choosing which vulnerabilities to patch**

- Assessment indexes:
  - How good in defense (Coverage)
  - How efficient (Efficiency)
  - How costly (Patching effort)

Exploit Prediction: Problem Definition **Remarks (I)**

- **MANY** factors **NOT** assessed

- Given a certain `vuln`:
  - How **many systems** do I have with `vuln`?
  - How **costly** is an **incident** based on `vuln`?
  - ...

Exploit Prediction: Problem Definition **Remarks (II)**

- **MANY** factors **NOT** assessed

- We observe which vulnerabilities have been **actually exploited worldwide** in T
- **Approximation** by collecting many intelligence feeds

- Given a certain `vuln`:
  - …
  - How **likely** is that **I** will be attacked with `vuln`?
  - …

Exploit Prediction: Example Criteria

Experiment Scenario

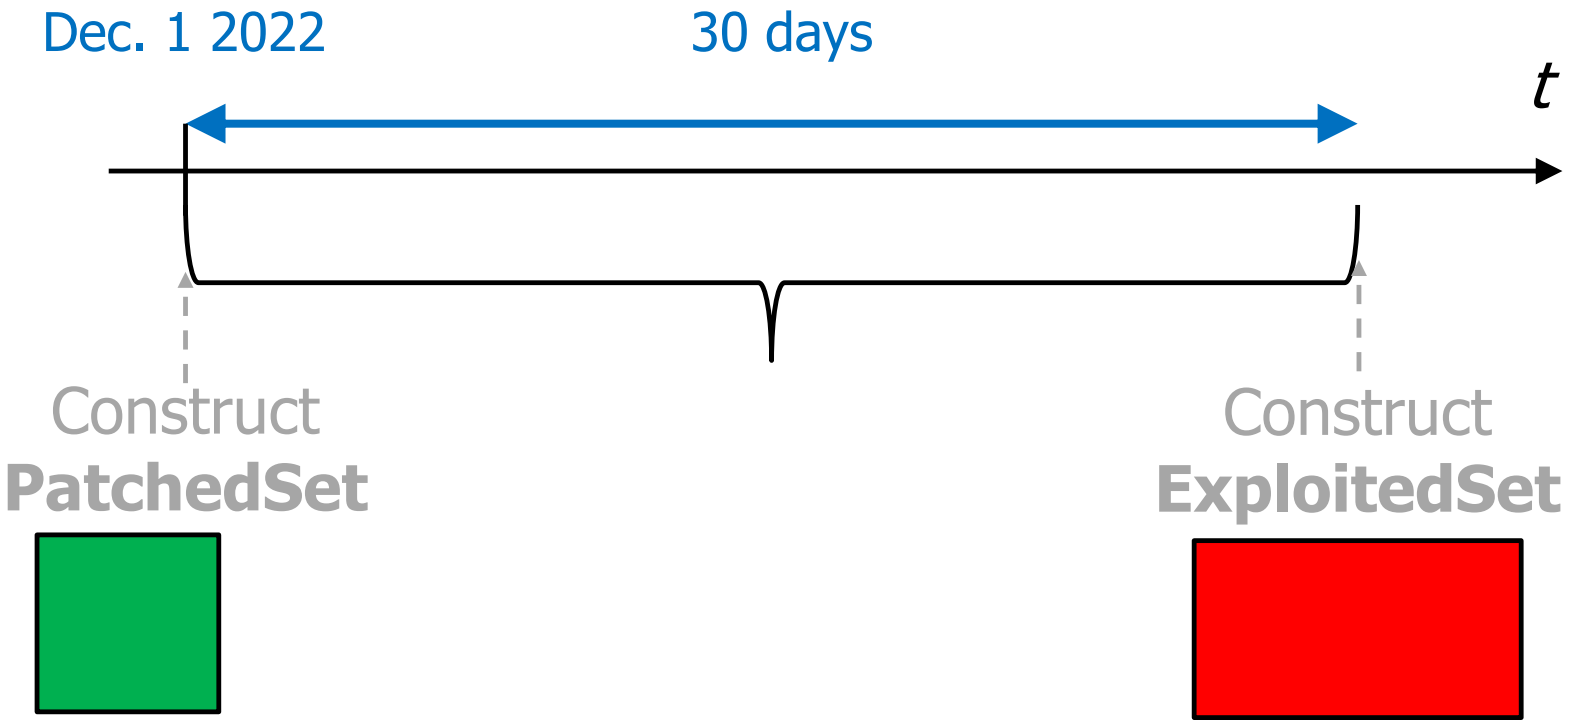

Various Criteria

≈8000 vulns observed

PatchedSet: CVSS (I)

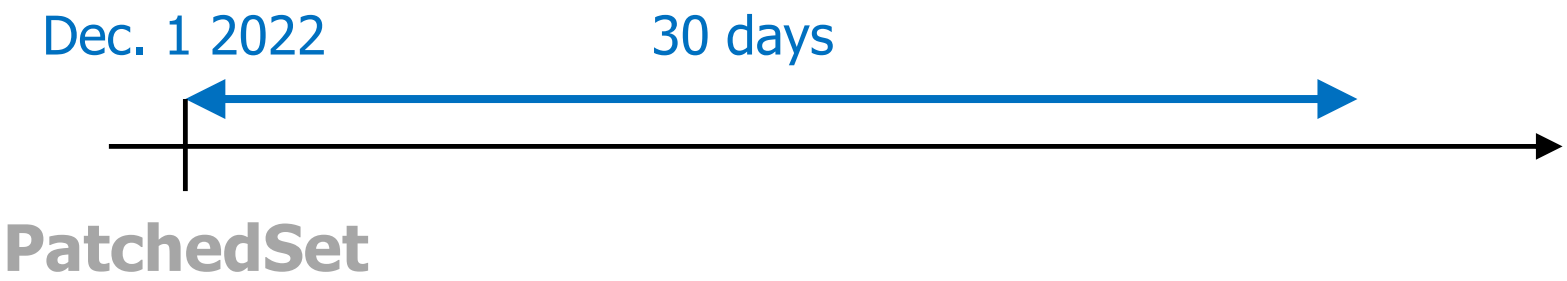

- All the CVEs with **CVSS >= 9.1** (≈15% of all vulns)
- Patching Effort: ≈28000 vulns

PatchedSet CVSS (II)

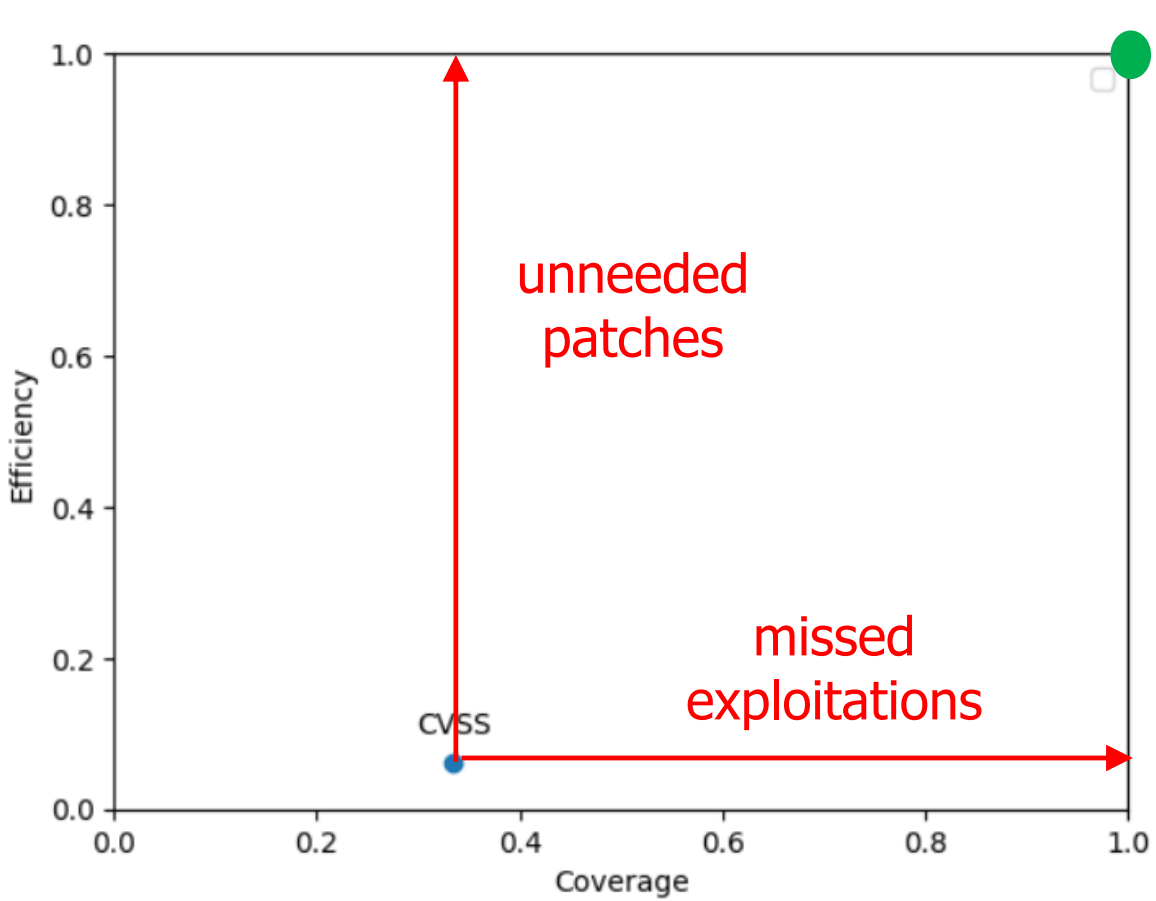

CVSS
is not
a good predictor of
exploitation

CISA - KEV

KNOWN EXPLOITED VULNERABILITIES CATALOG

PatchedSet: CISA-KEV (I)

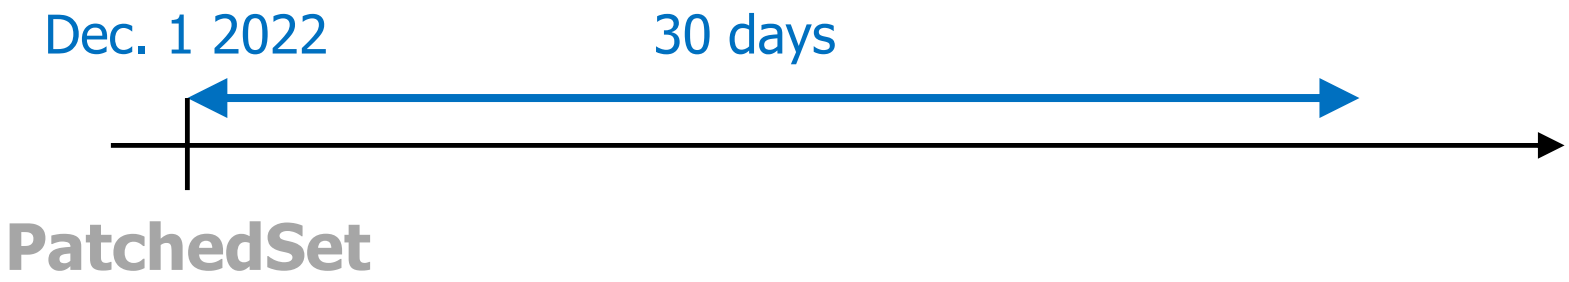

- CVEs **in CISA-KEV** (≈0.5% of all vulns)
- Patching Effort: ≈900 vulns

PatchedSet: CISA-KEV (II)

CISA-KEV
is not
a good predictor of exploitation

Selection based on Other Heuristics (I-a)

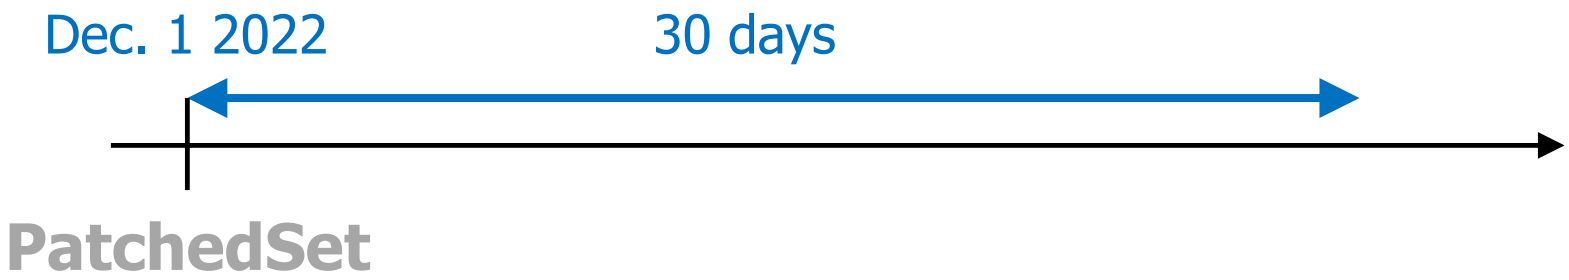

- Exploit in Metasploit
- Exploit in ExploitDB
- Impact is Code Execution
- Injection does not require Authentication

Selection based on Other Heuristics (I-b)

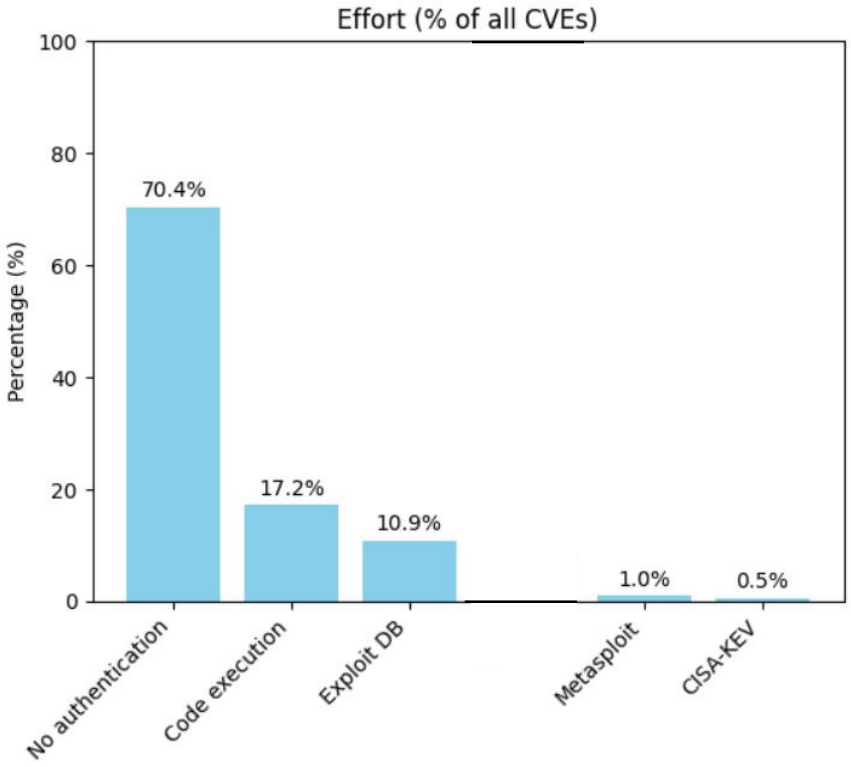

Selection based on Other Heuristics (II)

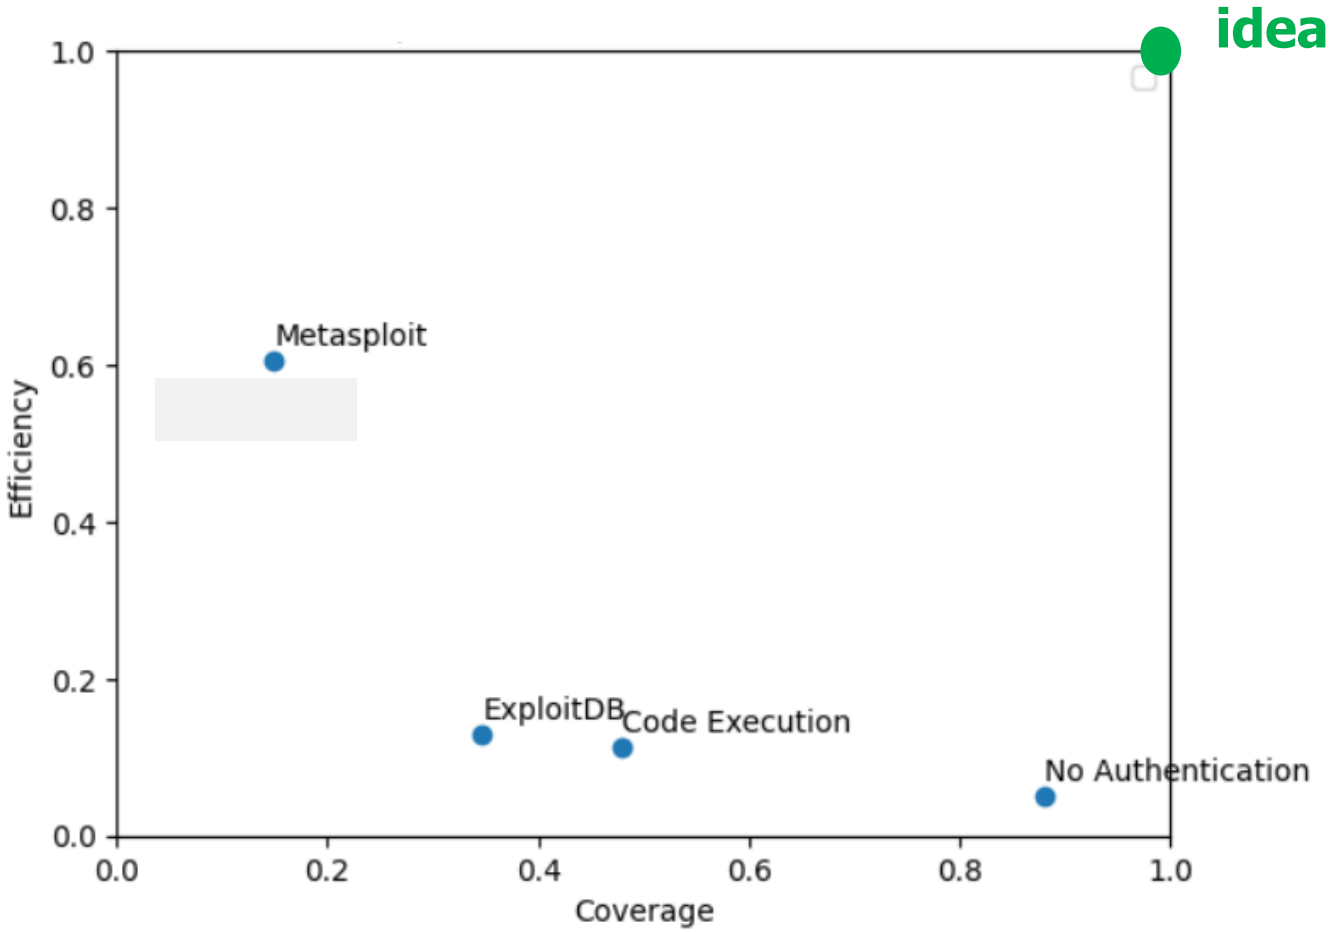

Exploit Prediction Scoring System (EPSS)

Exploit Probability: Definition (REMIND)

- Vulnerability `CVE-i`
- `P(CVE-i, d):`

  `# Exploitation attempts of CVE-i in [d,d+30]`
  `# Expolitation attempts of all CVEs in [d,d+30]`

- Probability that `CVE-i` will be exploited in the next 30 days

- **Worldwide** (everywhere)
- **Approximated** by collecting many TI feeds
- Computed **a posteriori**

EPSS Definition

- `EPSS(CVE-i, d)`: **Estimate** of `P(CVE-i, d)`

- Computed and published **daily** by the FIRST Consortium

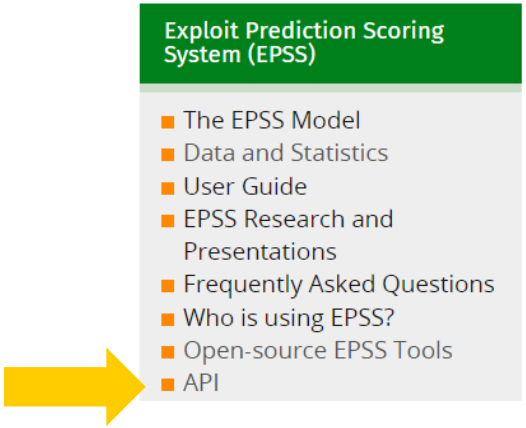

Enisa VD (I)


Home / Vulnerability List / EUVD-2024-49451
enisa
EUROPEAN UNION
VULNERABILITY
DATABASE
EUVD-2024-49451
Back to the vulnerability search
Severity
Alternative IDs
CVSS Base Score: 10 (v4.0)
(CVSS:4.0/AV:N/AC:L/AT:N/PR:N/UI:N/VC:H/VI:H/VA:H/SC:H/SI:H/SA:H)
CVE-2024-8878
EPSS Score
0.42
%
Summary
The password recovery mechanism for the forgotten password in Riello Netman 204 allows an attacker to reset the admin password and take over control of the device.This issue affects Netman 204: through 4.05.

Enisa VD (II)

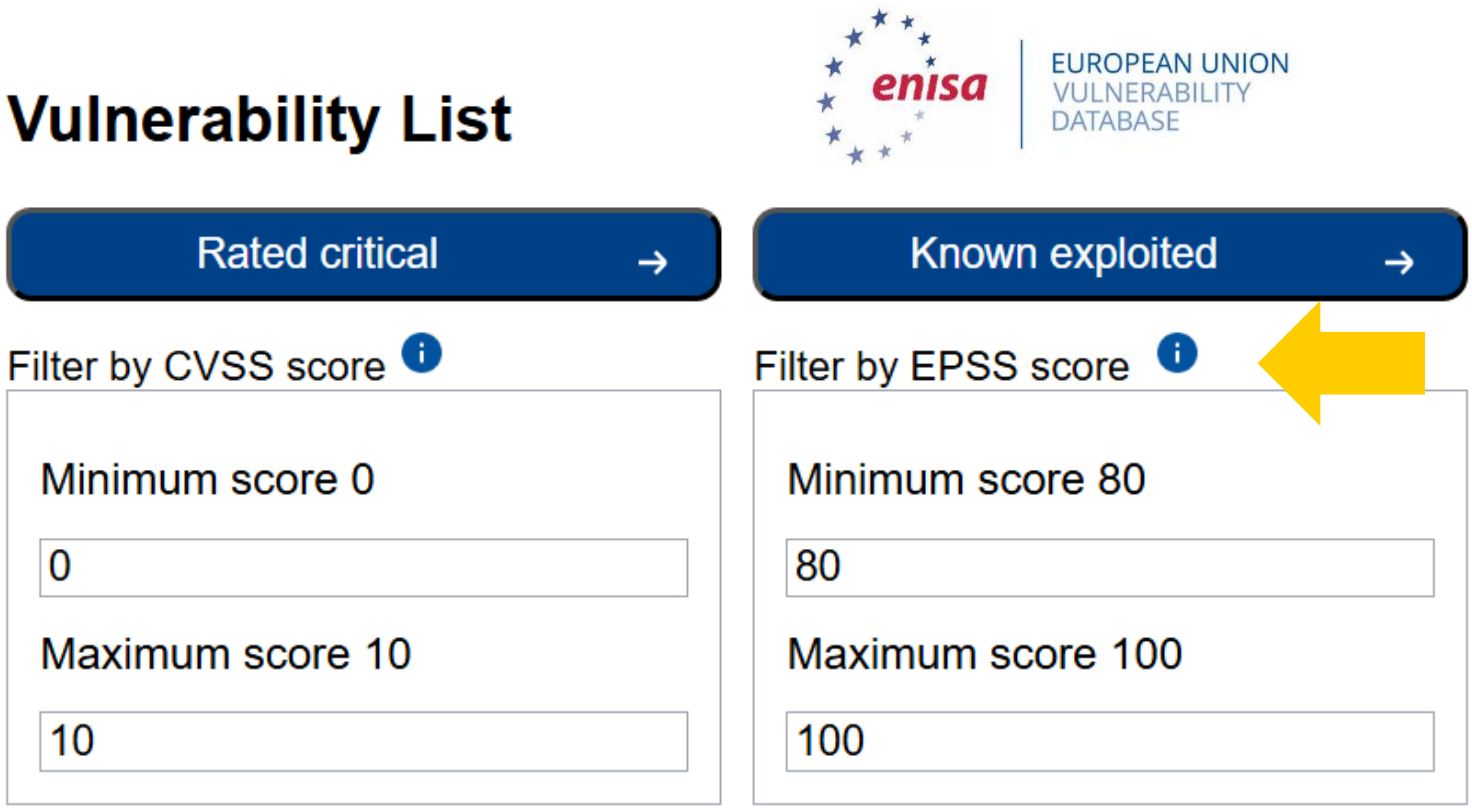
Vulnerability List
enisa
EUROPEAN UNION
VULNERABILITY
DATABASE
Rated critical
Known exploited
Filter by CVSS score
Filter by EPSS score
Minimum score 0
0
Maximum score 10
10
Minimum score 80
80
Maximum score 100
100

EPSS Computation

- Repeat **every day**:
  - For each CVE-i:
    - **Compute features** of CVE-i
    - **Estimate** its probability of exploitation in the next 30 days

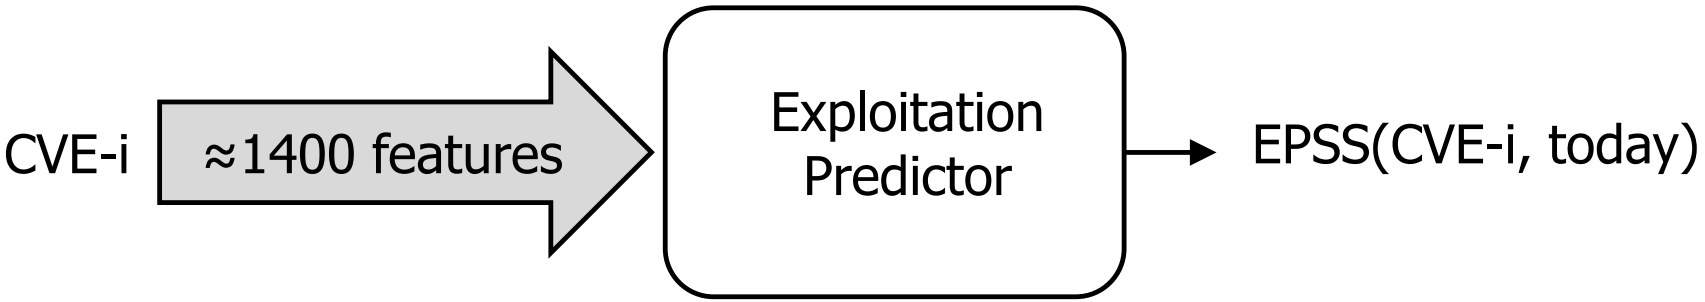

How is each CVE represented? (I)

- Array with 1400 elements
  - Numerical features
  - Categorical features (one-hot representation)

- Details out of scope
- **Information sources** in scope

How is each CVE represented? (II)

Description	Sources
Keyword description of vulnerability	Text description in MITRE CVE List
CVSS metrics	National Vulnerability Database (NVD)
CWE	National Vulnerability Database (NVD)
Vendor labels	National Vulnerability Database (NVD)
Age of the vulnerability	Days since CVE published in MITRE CVE list

Intrinsic properties of CVE-i
(more or less)

How is each CVE represented? (III)

Description	Sources
Exploitation activity in the wild (labels)	Fortinet, AlienVault, Shadowserver, GreyNoise
Publicly available exploit code	Exploit-DB, GitHub, MetaSploit
CVE mentioned on list or website	CISA KEV, Google Project Zero, Trend Micro ZDI
Social media	Mentions/discussion on Twitter
Offensive security tools and scanners	Intrigue, sn1per, jaeles, nuclei
References with labels	MITRE CVE List, NVD
Keyword description of vulnerability	Text description in MITRE CVE List
CVSS metrics	National Vulnerability Database (NVD)
CWE	National Vulnerability Database (NVD)
Vendor labels	National Vulnerability Database (NVD)
Age of the vulnerability	Days since CVE published in MITRE CVE list

- Summary of "**what people are saying** of CVE-i"
- Updated **daily**

CVE-I features → EPSS

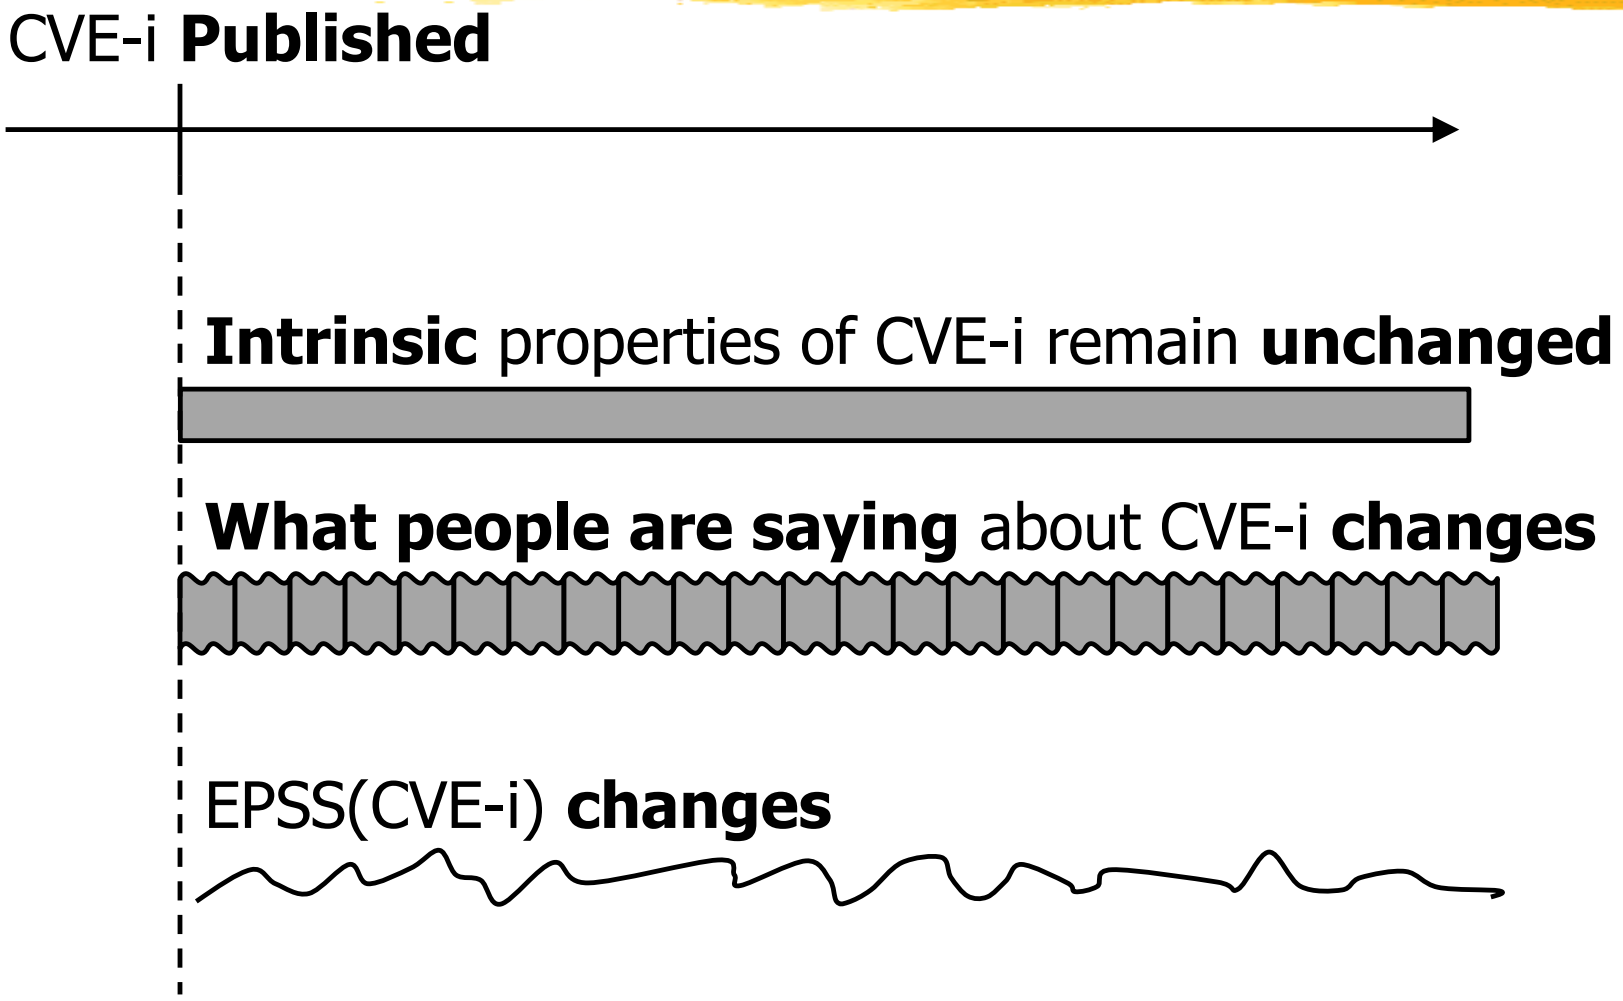

EPSS evolution: Examples (I)

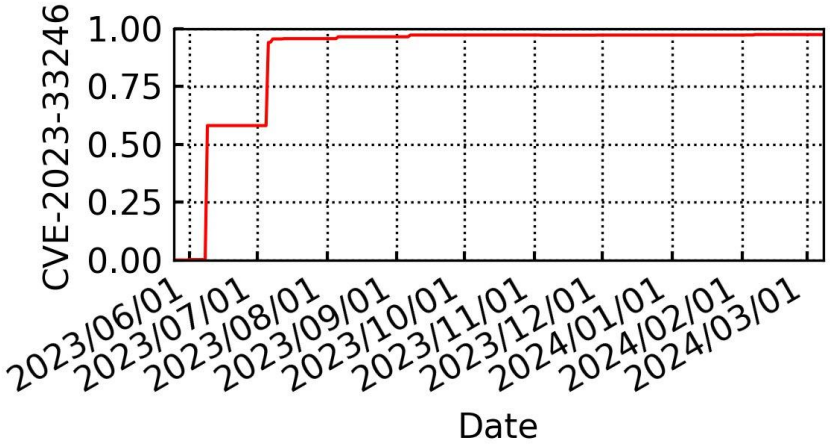

- Significant growth **after 1 day**
- ...and then again on the next day
- Heavily exploited for more than 6 months

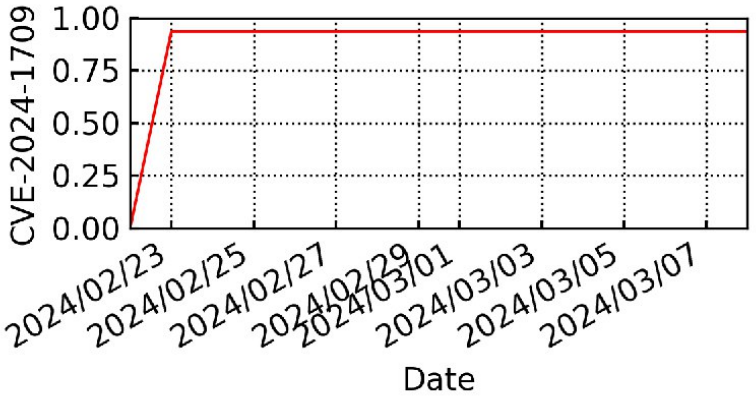

- **Immediately** exploited heavily, for several weeks

EPSS evolution: Examples (II)

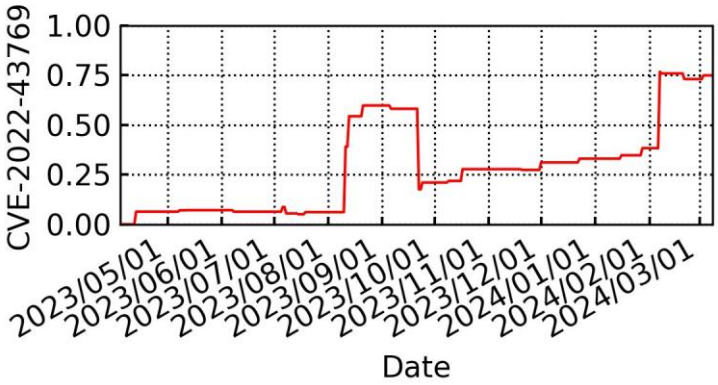

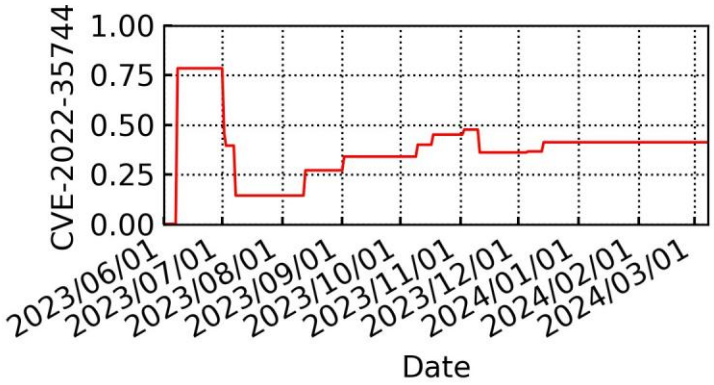

- Temporal evolution may often be:
  - Very "irregular"
  - Very hard to predict (even in the short term)

How is EPSS computed?

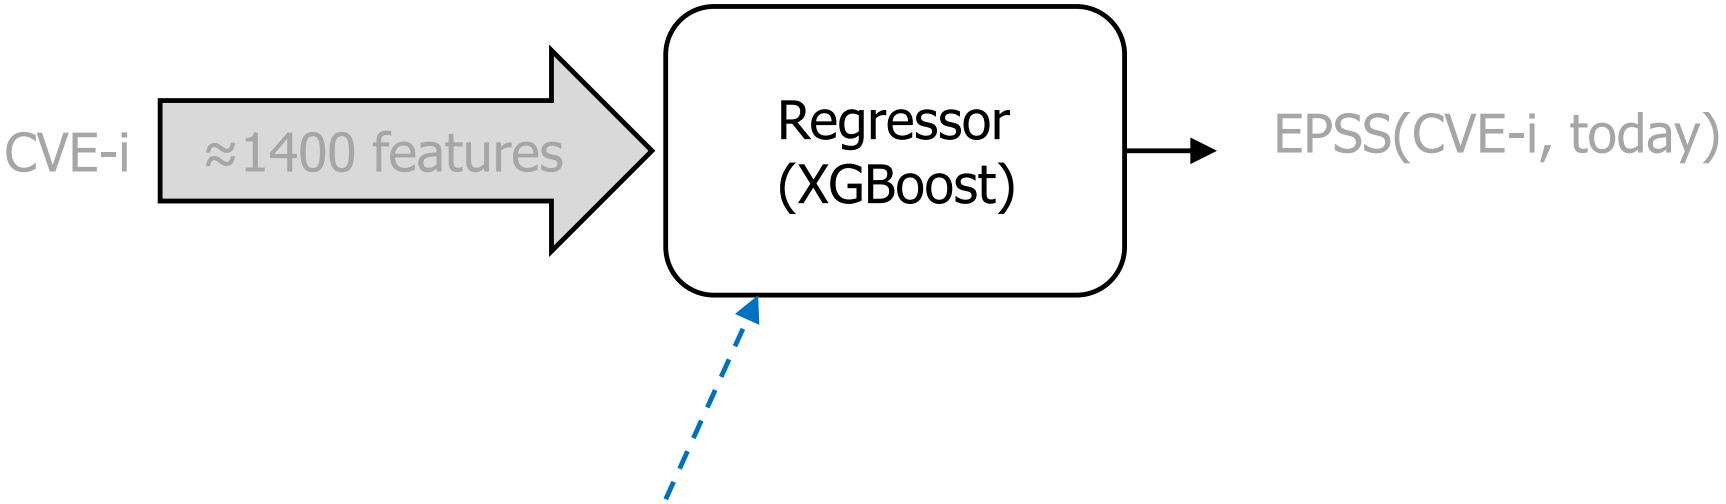

- Data driven model
- Trained on 1 year of data

- March 2023: 3rd model refinement
- March 2025: 4th release

P_observed(CVE-i) vs P_predicted(CVE-i)

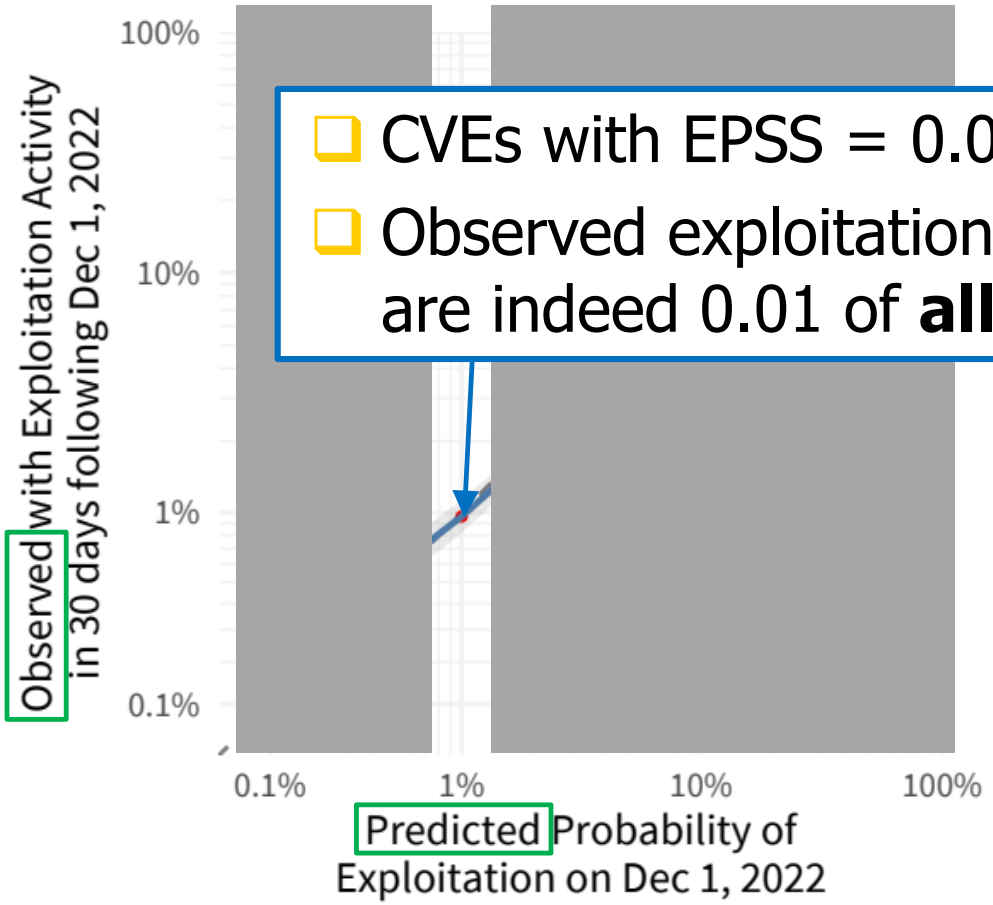

- CVEs with EPSS = 0.01
- Observed exploitation attempts of these CVEs are indeed 0.01 of **all** the observed attempts

P_observed(CVE-i) ≈ P_predicted(CVE-i)

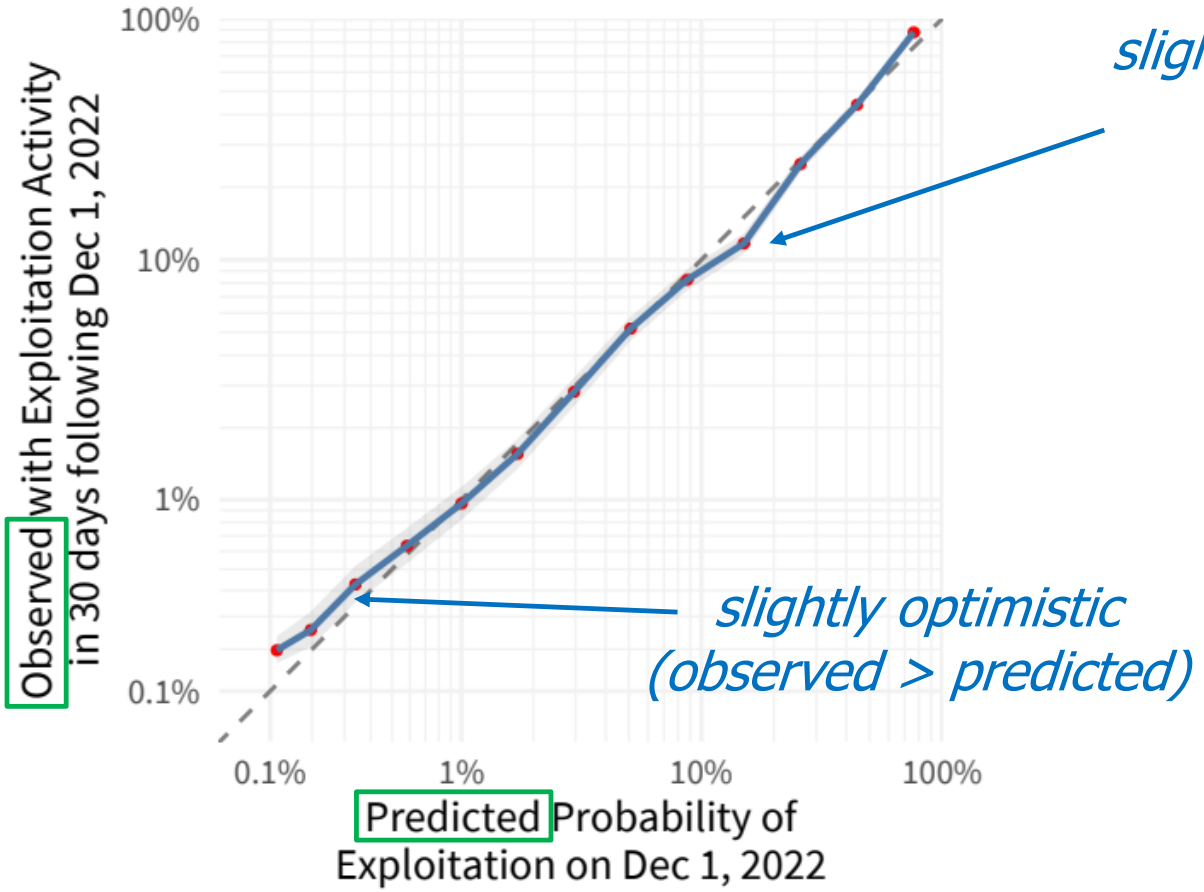

EPSS for Exploit Prediction: Coverage and Efficiency?

EPSS vs CVSS

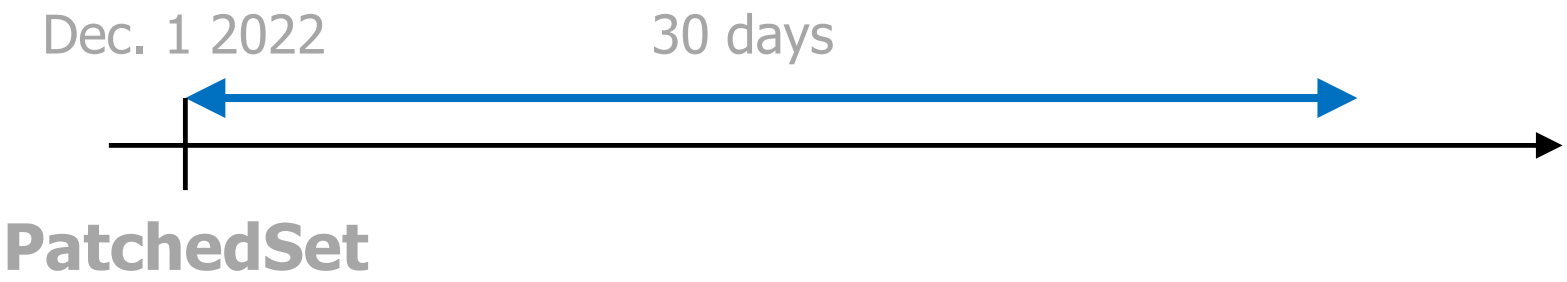

- CVEs with "large" **CVSS**
- CVEs with "large" **EPSS**

- Which criterion is better?
- Is it better to prioritize patches based on CVSS…or on EPSS?

EPSS vs CVSS: Identical Patching Effort (I)

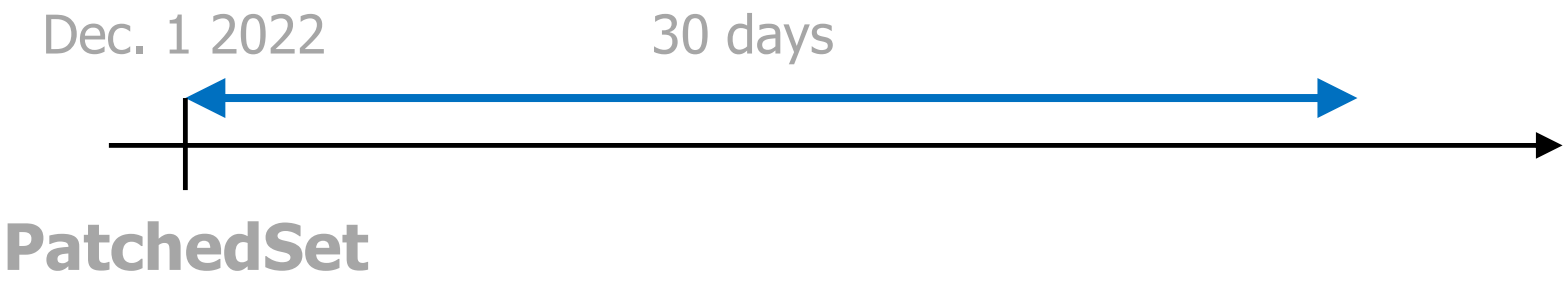

- Set1: CVSS(CVE-i) >= 9.1 (**15% of all CVEs**)
- Set2: EPSS(CVE-i) >= **0.022** (**15% of all CVEs**)

- **Identical Patching Effort**
- Efficiency?
- Coverage?

EPSS vs CVSS: Identical Patching Effort (II)

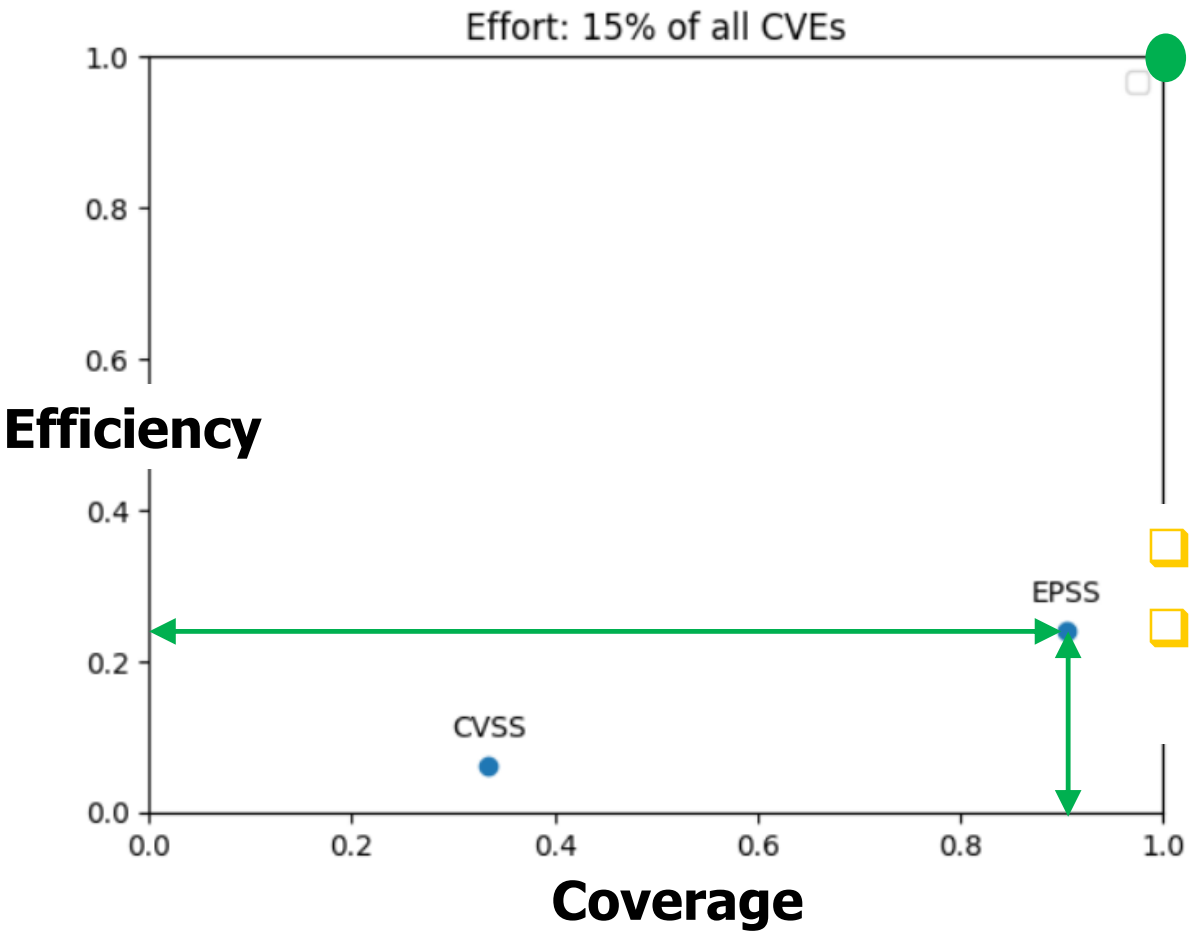

- Excellent coverage
- More efficient than CVSS (But not very efficient)

EPSS vs CVSS: Identical Coverage (I)

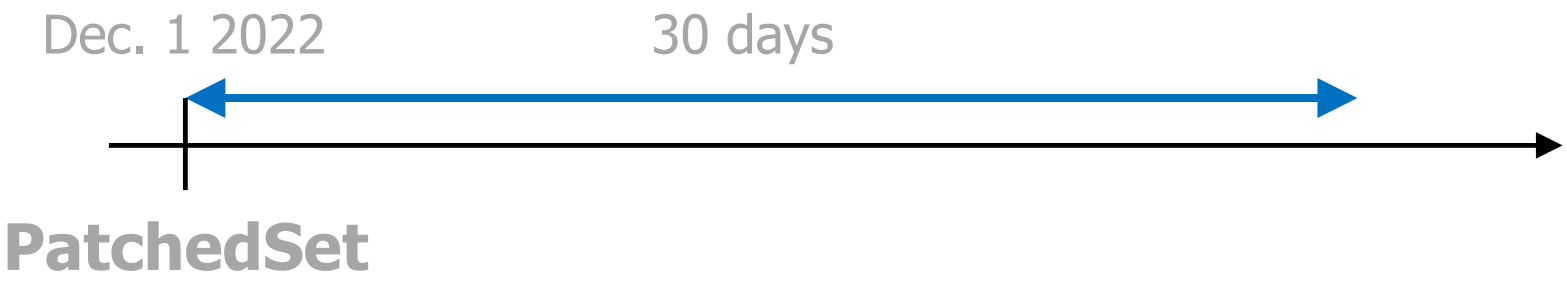

- Set1: CVSS(CVE-i) >= 7 (**Coverage 82%**)
- Set2: EPSS(CVE-i) >= **0.088** (**Coverage 82%**)

- **Identical Coverage**
- Efficiency?
- Patching Effort?

EPSS vs CVSS: Identical Coverage (II-a)

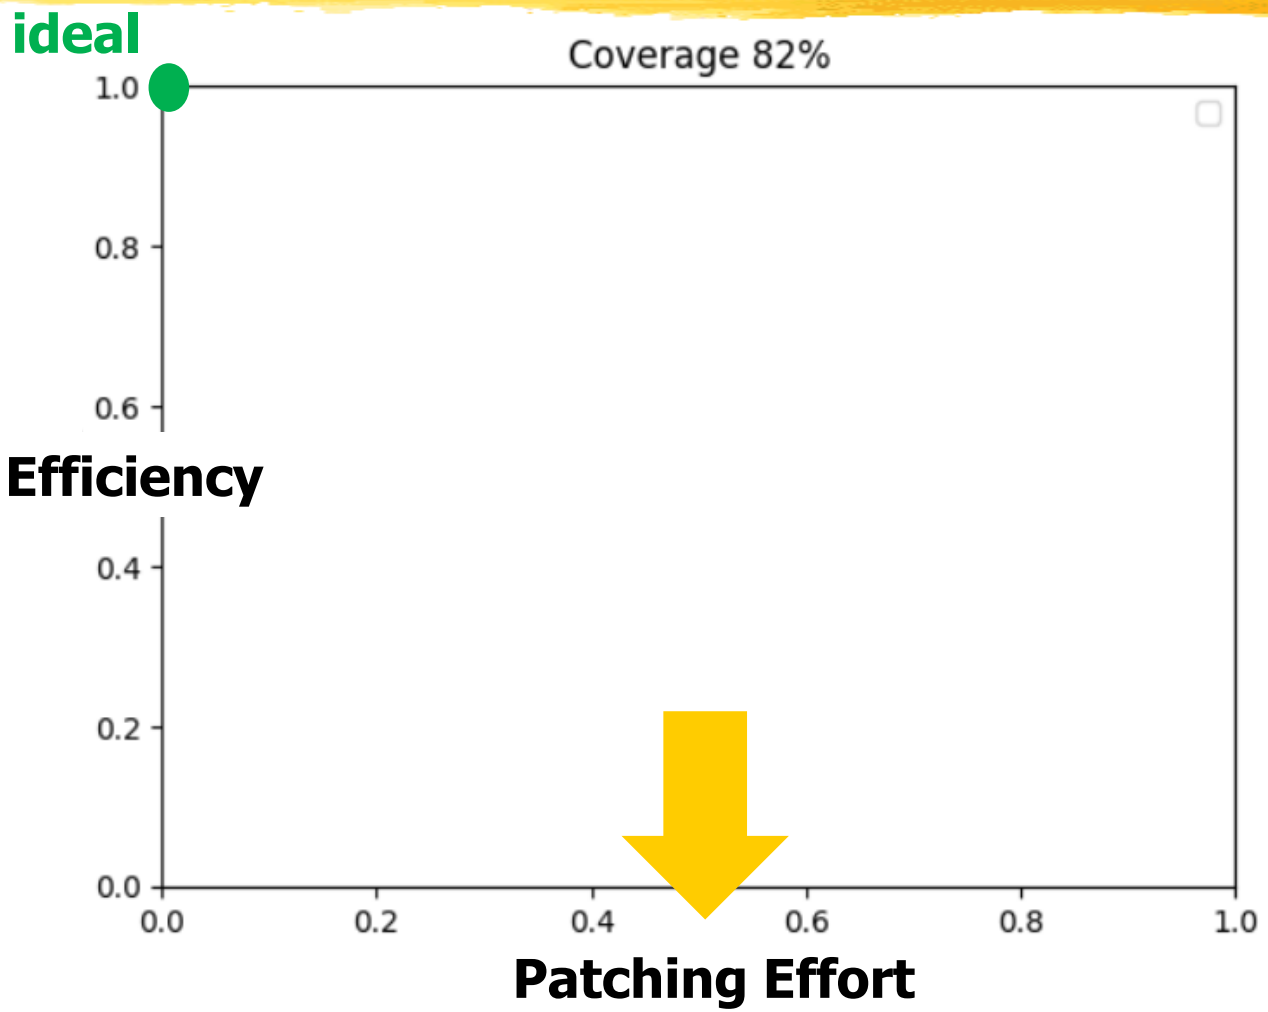

EPSS vs CVSS: Identical Coverage (II-b)

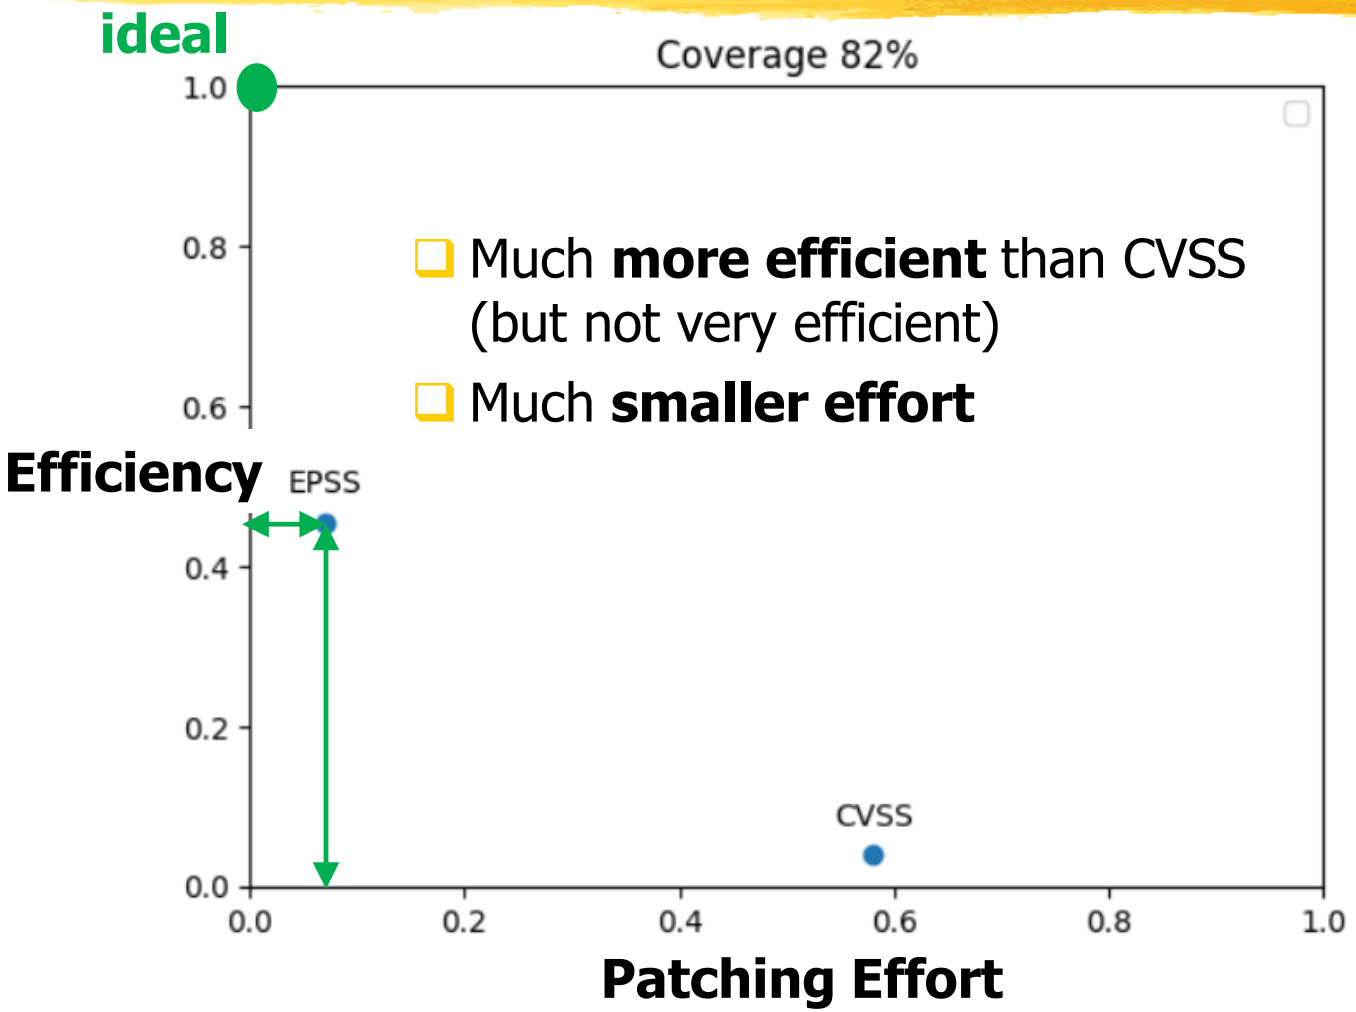

- Much **more efficient** than CVSS (but not very efficient)
- Much **smaller effort**

EPSS vs Other Heuristics (I)

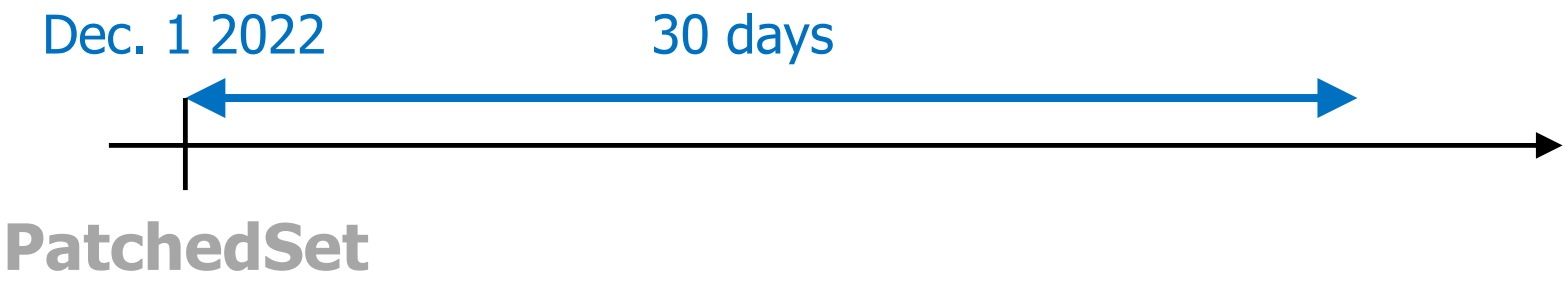

- Exploit in Metasploit
- Exploit in ExploitDB
- Impact is Code Execution
- Injection does not require Authentication

- EPSS > 8.8%

EPSS vs Other Heuristics (II-a)

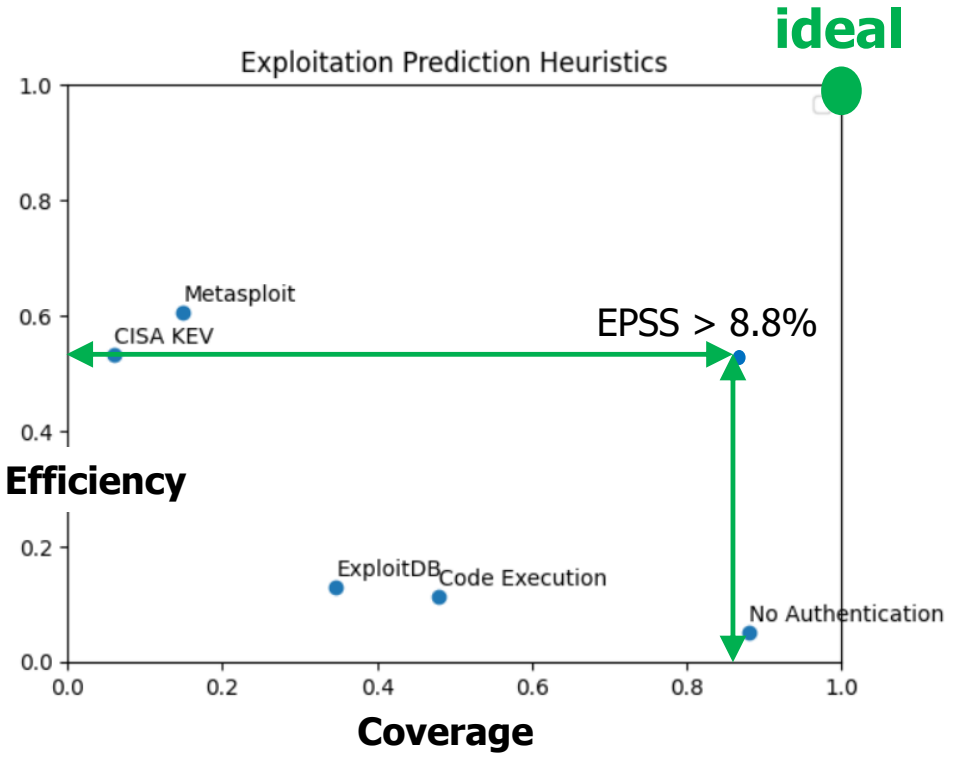

EPSS vs Other Heuristics (II-a)

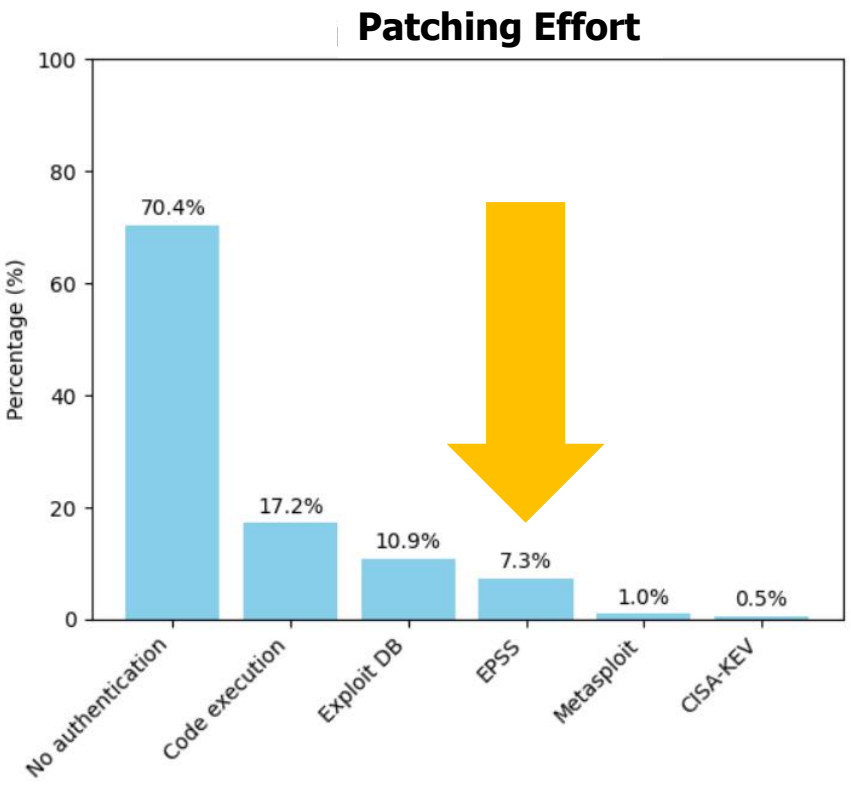

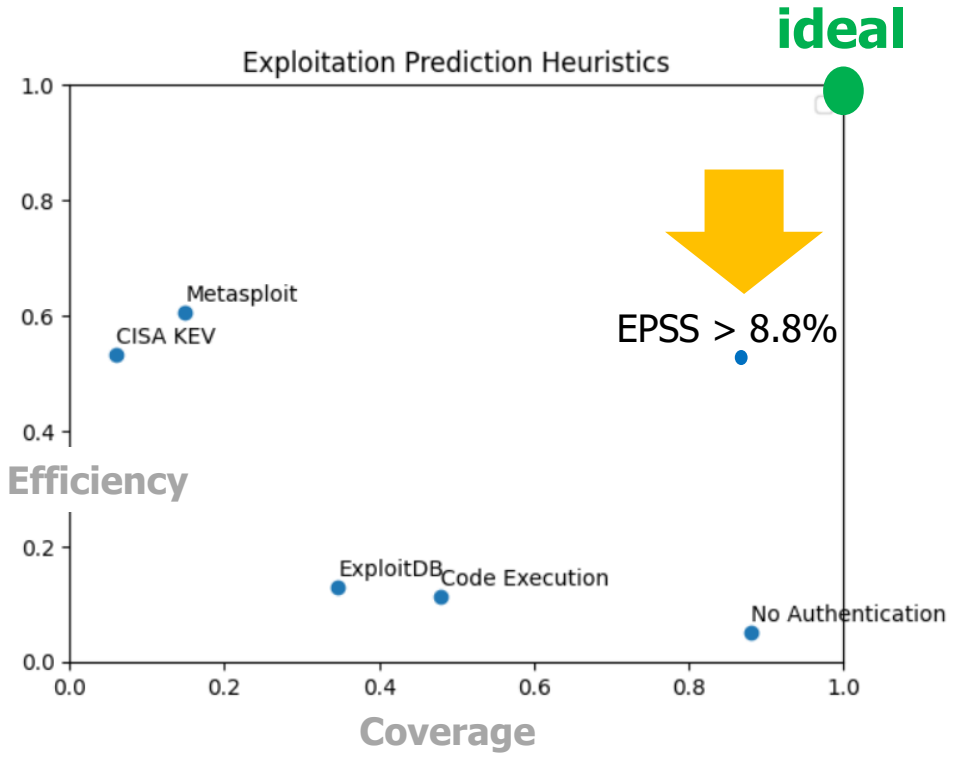

Understanding EPSS for Exploit Prediction

- Based on some of our (yet unpublished) analyses

Using EPSS in practice

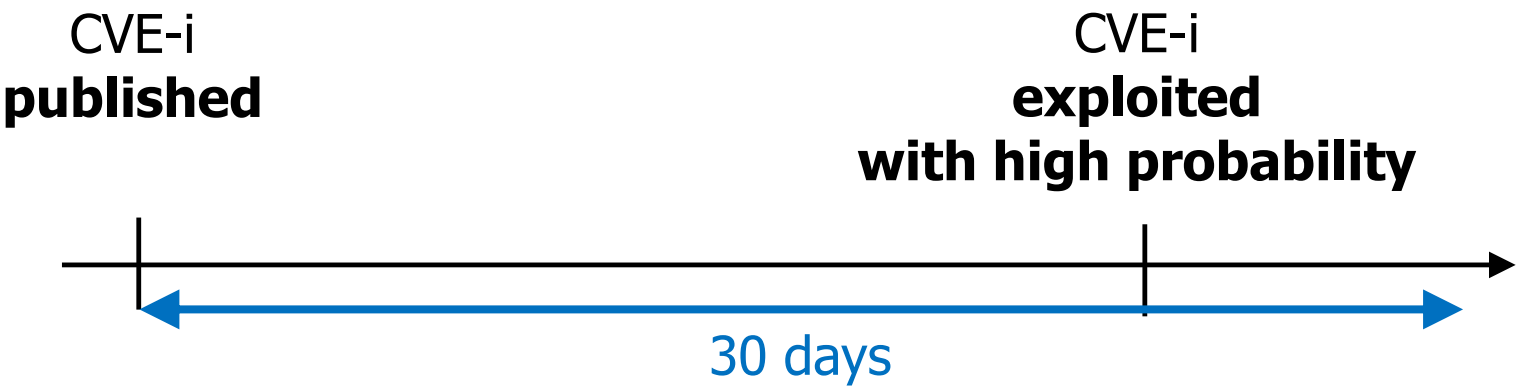

- Prioritizing CVE-i with "large" EPSS is a good criterion

- But what about **newly published** CVE-i?
- **When** we will prioritize them?

Does EPSS predict FUTURE exploitation?

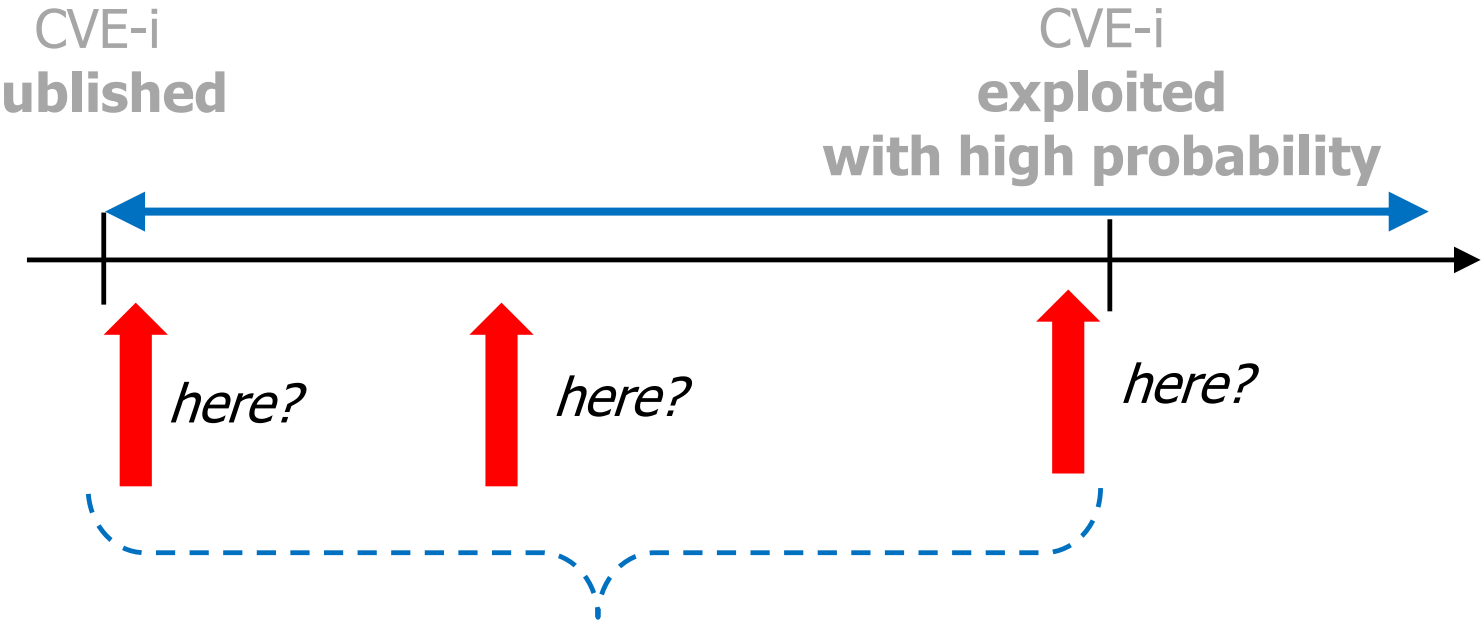

Day on which EPSS(CVE-i) becomes "large enough"

...so that CVE-i is prioritized

...or does it summarize what is ALREADY happening?

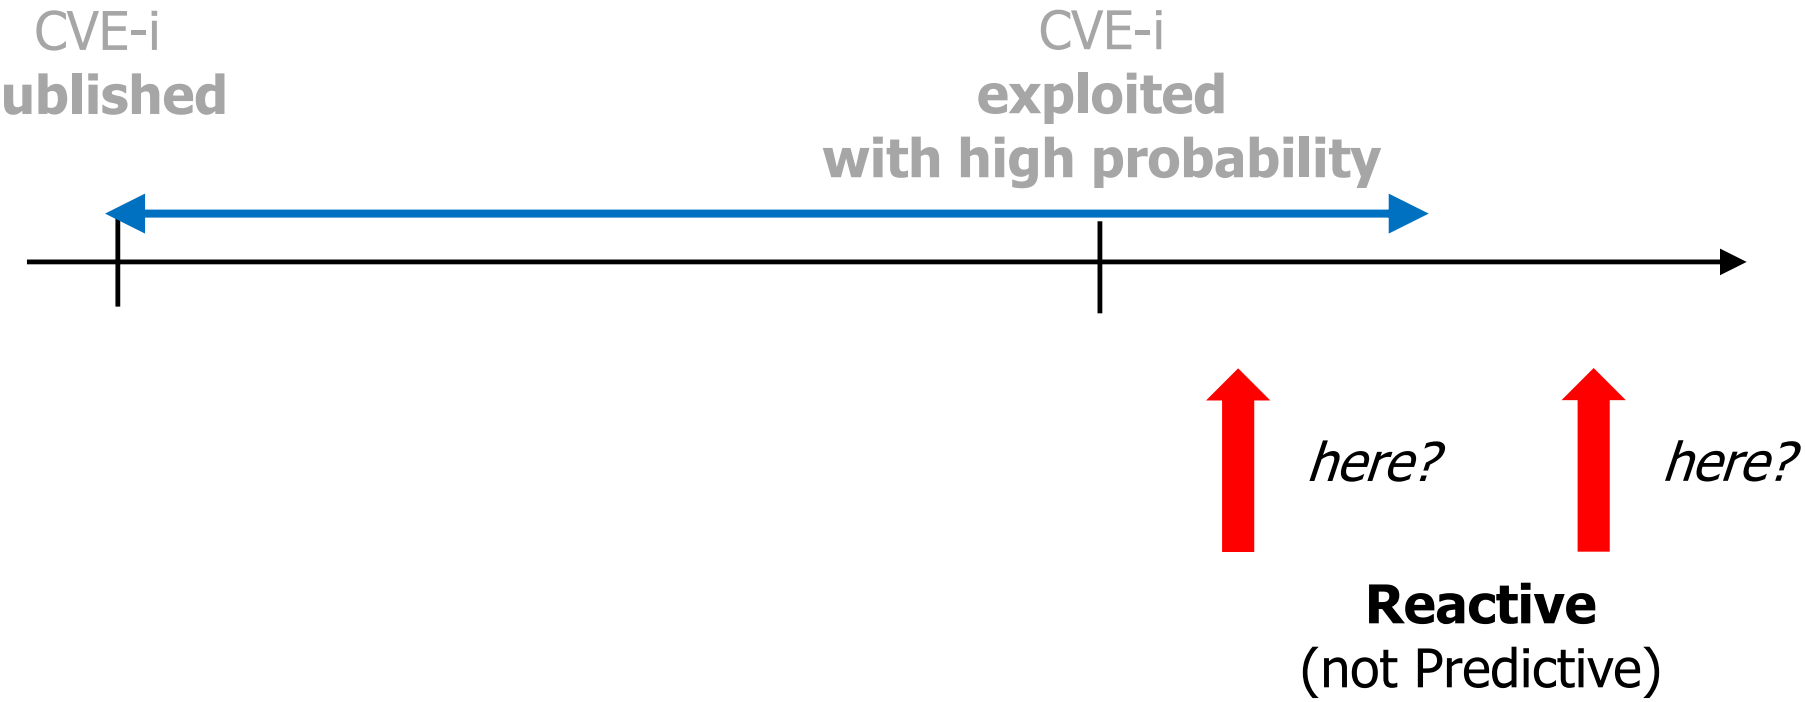

EPSS: Dynamic behavior (I)

- Prioritized-CVE :=
  - All CVE-i **published** in [March 2023, September 2024]
    $\rightarrow$ 45080
  - Such that **EPSS(CVE-i) > 0.7** (at some day)
    $\rightarrow$ 137

- Does EPSS
  - "**predict** actual exploitation by **many** days"?
  - "**predict** actual exploitation by **a few** days"?
  - "**summarizes after** the fact what is **already** happening"?

EPSS: Dynamic behavior (II)

- Does EPSS
  - "**predict** actual exploitation by **many** days"?
  - "**predict** actual exploitation by **a few** days"?
  - "**summarizes after** the fact what is **already** happening"?
- We **cannot** answer

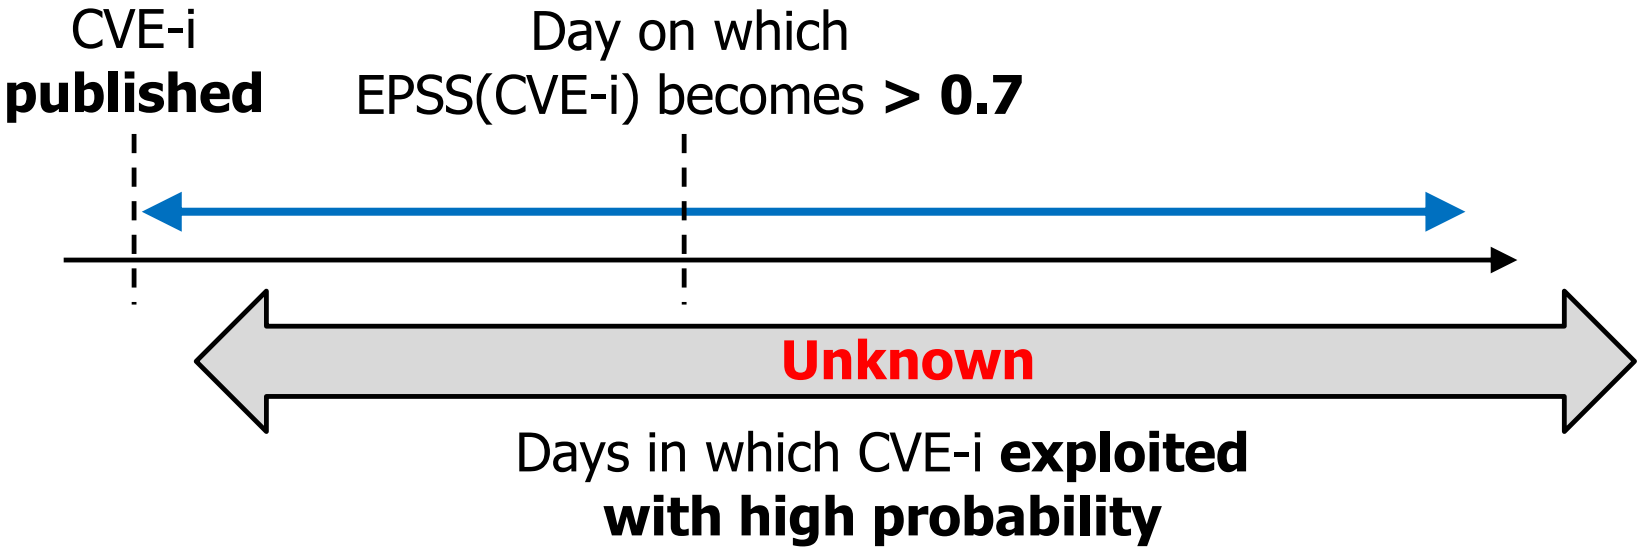

EPSS: Dynamic behavior (III)

- Does EPSS
  - "**predict** actual exploitation by **many** days"?
  - "**predict** actual exploitation by **a few** days"?
  - "**summarizes after** the fact what is **already** happening"?
- We **cannot** answer

- We answer a **different** question:
  How long does it take for the EPSS to reach 0.7?
  - A few **days**?
  - …or a few **weeks**?
  - …or a few **months**?

EPSS: Dynamic behavior (II)

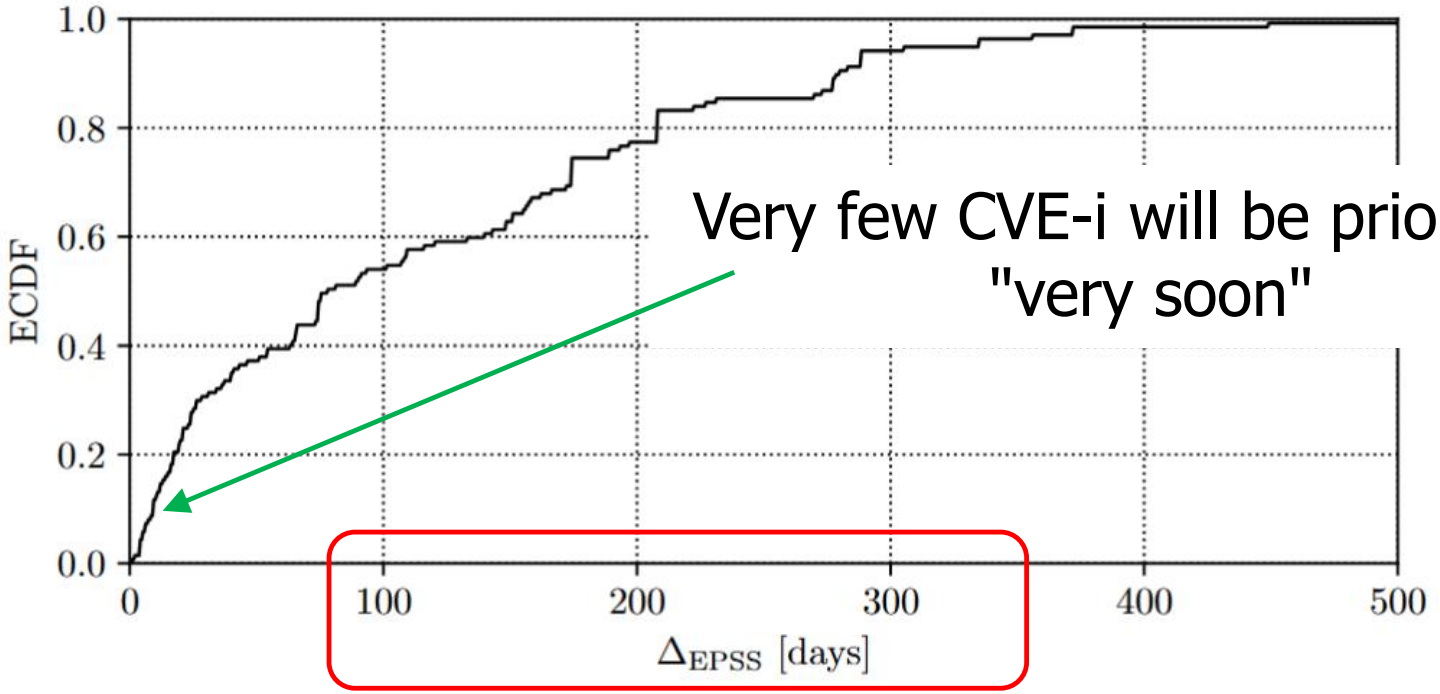

Fig. 5. Distribution of the time taken to reach the EPSS threshold for the vulnerabilities in NVD_{EPSS}

CVEs whose EPSS becomes > 0.7

What is the influence of CISA KEV? (I)

Description	Sources	
Exploitation activity in the wild (labels)	Fortinet, AlienVault, Shadowserver, GreyNoise	
Publicly available exploit code	Exploit-DB, GitHub, MetaSploit	
CVE mentioned on list or website	CISA KEV, Google Project Zero, Trend Micro ZDI	
Social media	Mentions/discussion on Twitter	
Offensive security tools and scanners	Intrigue, sn1per, jaeles, nuclei	**TI**
References with labels	MITRE CVE List, NVD	**INTRINSIC**
Keyword description of vulnerability	Text description in MITRE CVE List	
CVSS metrics	National Vulnerability Database (NVD)	
CWE	National Vulnerability Database (NVD)	
Vendor labels	National Vulnerability Database (NVD)	
Age of the vulnerability	Days since CVE published in MITRE CVE list	

- How likely that EPSS(CVE-i) becomes > 0.7 **before inclusion** in CISA KEV ?
- IF not very likely
  THEN probably it is more reactive than predictive

What is the influence of CISA KEV? (II)

Very few CVE-i will be prioritized **before** being inserted in KEV

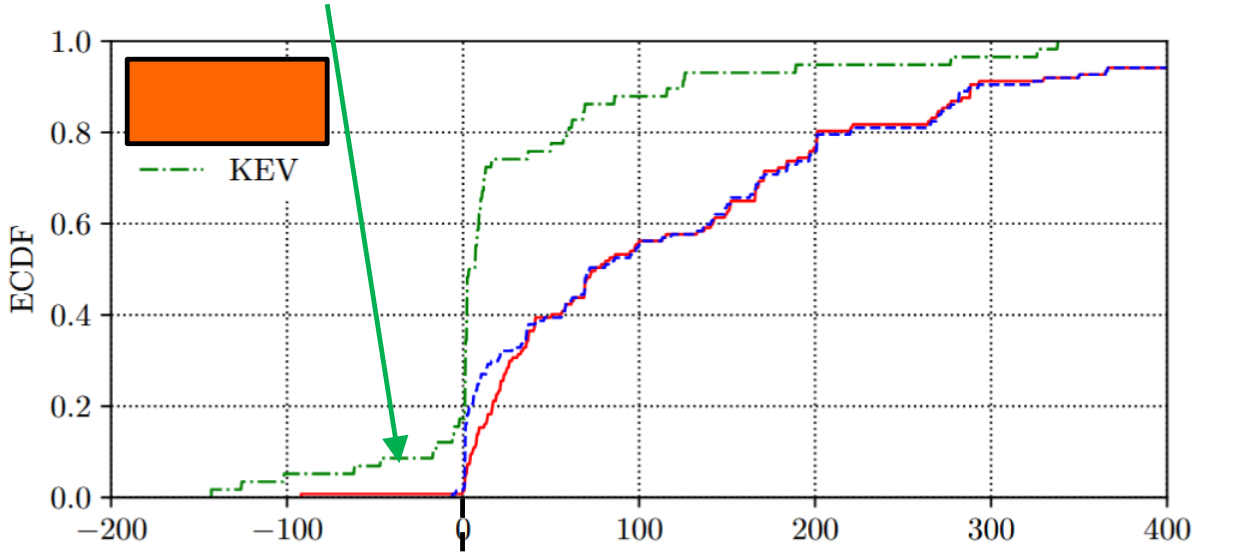

EPSS becomes > 0.7 ***before*** *inclusion in KEV*

EPSS becomes >0.7 ***after*** *inclusion in KEV*

Remark

- E**P**SS(CVE-i, d):
  - Probability **estimate** that CVE-i **will be** exploited in [d, d+30]
- Updated daily
- Based on:
  - Intrinsic properties of CVE-I
  - What people are saying about CVE-i

- It is called a **predictor**
- …but there are strong indications that it is more a **summary** of what is **already** happening